

# **ANNAI MIRA COLLEGE OF ENGINEERING AND TECHNOLOGY**

NH-46, Chennai-Bengaluru National Highways, Arappakkam,

Ranipet-632517, Tamil Nadu, India

Telephone: 04172-292925 Fax: 04172-292926

Email:[amcet.rtet@gmail.com](mailto:amcet.rtet@gmail.com)/[info@amcet.in](mailto:info@amcet.in) Web: [www.amcet.in](http://www.amcet.in)

## **DEPARTMENT OF INFORMATION TECHNOLOGY**



### **CCS334 – BIG DATA ANALYTICS LABORATORY**

**Name** : .....

**Register Number** : .....

**Year & Branch** : .....

**Semester** : .....

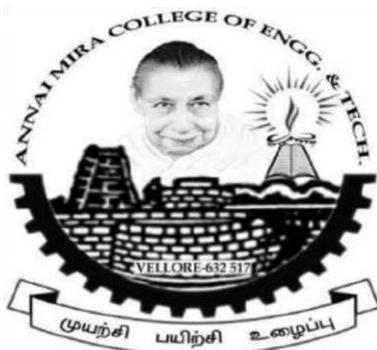
**Academic Year** : .....

# **ANNAI MIRA COLLEGE OF ENGINEERING AND TECHNOLOGY**

NH-46, Chennai-Bengaluru National Highways, Arappakkam,

Ranipet-632517, TamilNadu, India

Telephone: 04172-292925 Fax: 04172-292926



## **CERTIFICATE**

This is to Certify that the Bonafide record of the practical work done by..... Register Number .....of III<sup>rd</sup> year B.Tech (Information Technology) submitted for the B.Tech-Degree practical examination (Vth Semester) in **CCS334 BIG DATA ANALYTICS LABORATORY** during the academic year **2025-2026**.

**Staff in –Charge**

**Head of the Department**

Submitted for the practical examination held on -----

**Internal Examiner**

**External Examiner**

## **TABLE OF CONTENTS**

| <b>EXP.<br/>NO</b> | <b>DATE</b> | <b>LIST OF EXPERIMENTS</b>  | <b>PAGE<br/>NO</b> | <b>SIGNATURE</b> |
|--------------------|-------------|---|--------------------|------------------|
| 1                  |             | Downloading and Installing Hadoop; Understanding Different Hadoop Modes. Startup Scripts, Configuration Files.            | 1                  |                  |
| 2                  |             | Hadoop Implementation of File Management Tasks, Such as Adding Files and Directories, Retrieving Files and Deleting Files | 13                 |                  |
| 3                  |             | Implement of Matrix Multiplication with Hadoop Map Reduce   | 17                 |                  |
| 4                  |             | Run a Basic Word Count Map Reduce Program to Understand Map Reduce Paradigm.  | 22                 |                  |
| 5                  |             | Installation of Hive along with Practice Examples.  | 28                 |                  |
| 6                  |             | Installation of Hbase along with Practice Examples  | 32                 |                  |
| 7                  |             | Installation of Thrift.   | 36                 |                  |
| 8                  |             | Practice Importing and Exporting Data From Various Databases.   | 38                 |                  |

|                 |   |
|-----------------|---|
| <b>EX. No:1</b> | <b>DOWNLOADING AND INSTALLING HADOOP<br/>UNDERSTANDING DIFFERENT HADOOP MODES.<br/>STARTUPSCRIPTS, CONFIGURATION FILES.</b> |
| <b>DATE:</b>    |   |

### **AIM:**

To Download and installing hadoop software to create different hadoop modes.

### REREQUISITES TO INSTALL HADOOP ON WINDOWS

- **VIRTUAL BOX** (For Linux): it is used for installing the operating system on it.
- **OPERATING SYSTEM:** You can install Hadoop on Windows or Linux based operating systems. Ubuntu and CentOS are very commonly used.
- **JAVA:** You need to install the Java 8 package on your system.
- **HADOOP:** You require Hadoop latest version

#### **1. Install Java**

- Java JDK Link to download <https://www.oracle.com/java/technologies/javase-jdk8-downloads.html>
- Extract and install Java in C:\Java
- Open cmd and type -> javac –version

#### **Command Prompt**

```
Microsoft Windows [Version 10.0.19041.572]
(c) 2020 Microsoft Corporation. All rights reserved.

C:\Users\asus>javac -version
javac 1.8.0_241
```

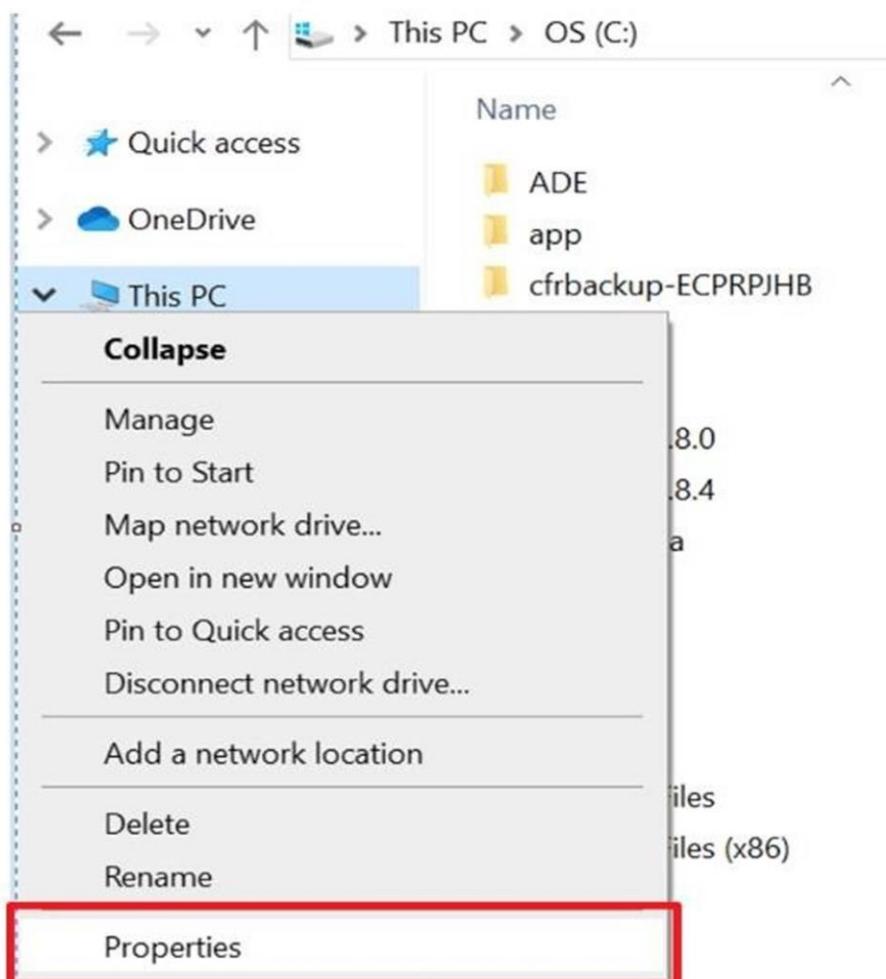
#### **2. Download Hadoop**

<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.3.0/hadoop-3.3.0.tar.gz>

extract to C:\Hadoop

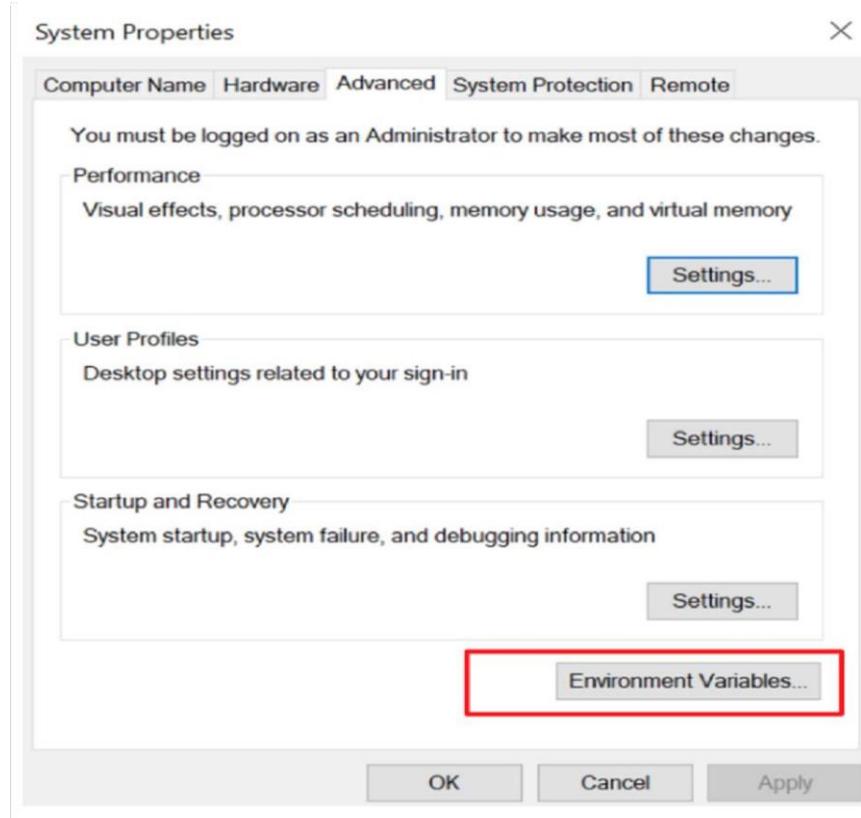
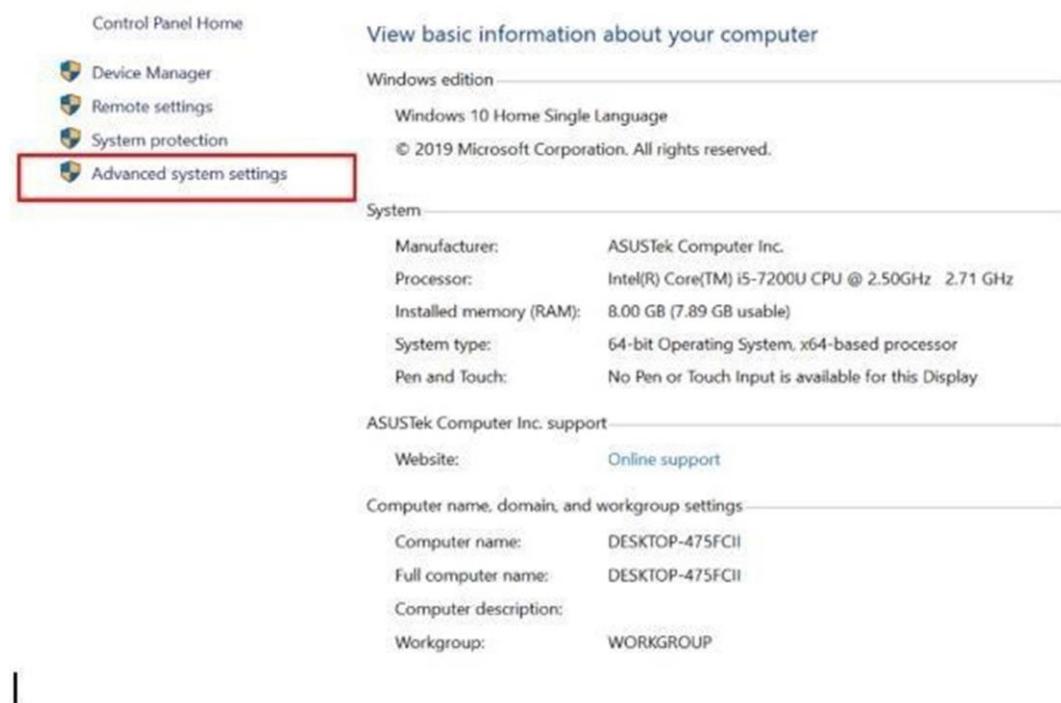
|                    |                    |             |
|--------------------|--------------------|-------------|
| ADE                | 1/26/2020 11:13 AM | File folder |
| app                | 1/26/2020 10:53 AM | File folder |
| cfrbackup-ECPRPJHB | 4/18/2019 10:25 PM | File folder |
| eSupport           | 7/13/2017 5:22 AM  | File folder |
| Games              | 8/20/2019 9:40 PM  | File folder |
| hadoop             | 11/8/2020 3:15 PM  | File folder |
| hadoop-2.8.0       | 12/10/2019 3:02 PM | File folder |
| hadoop-2.8.4       | 6/14/2019 9:36 PM  | File folder |
| hadoop-3.3.0       | 11/8/2020 4:30 PM  | File folder |
| Hortonwork         | 11/8/2020 2:40 PM  | File folder |
| Informatica        | 1/28/2020 12:52 AM | File folder |
| Java               | 11/8/2020 3:25 PM  | File folder |
| logs               | 3/27/2020 9:36 PM  | File folder |
| oraclexe           | 1/29/2020 11:52 PM | File folder |

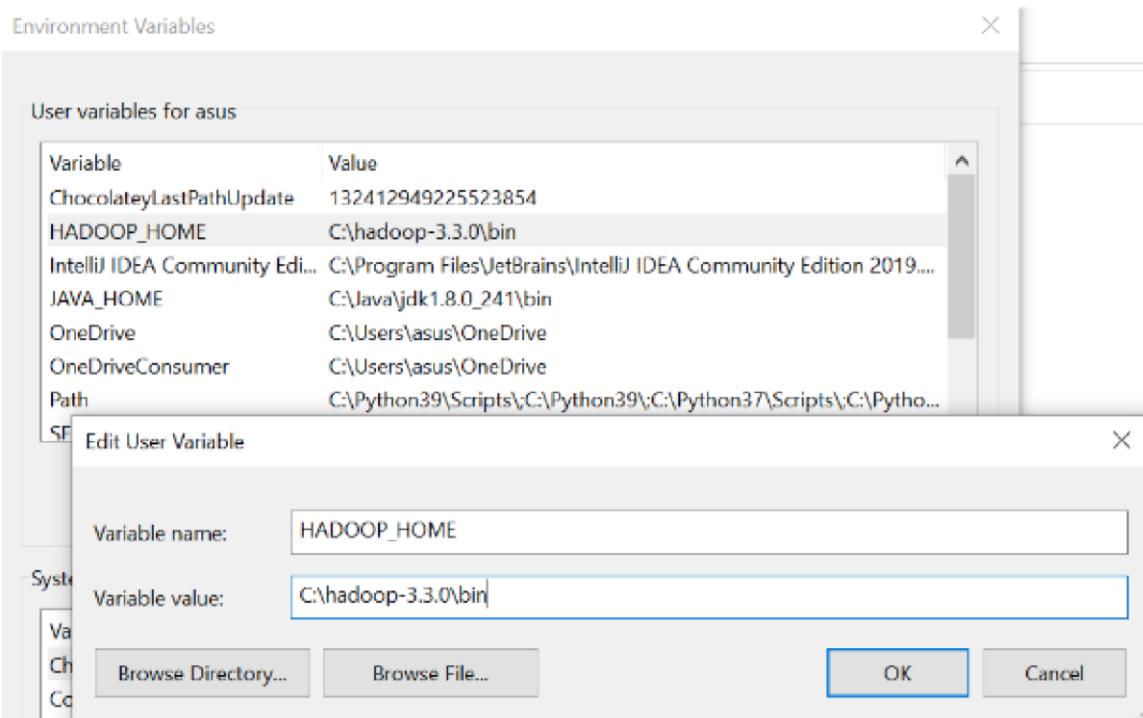
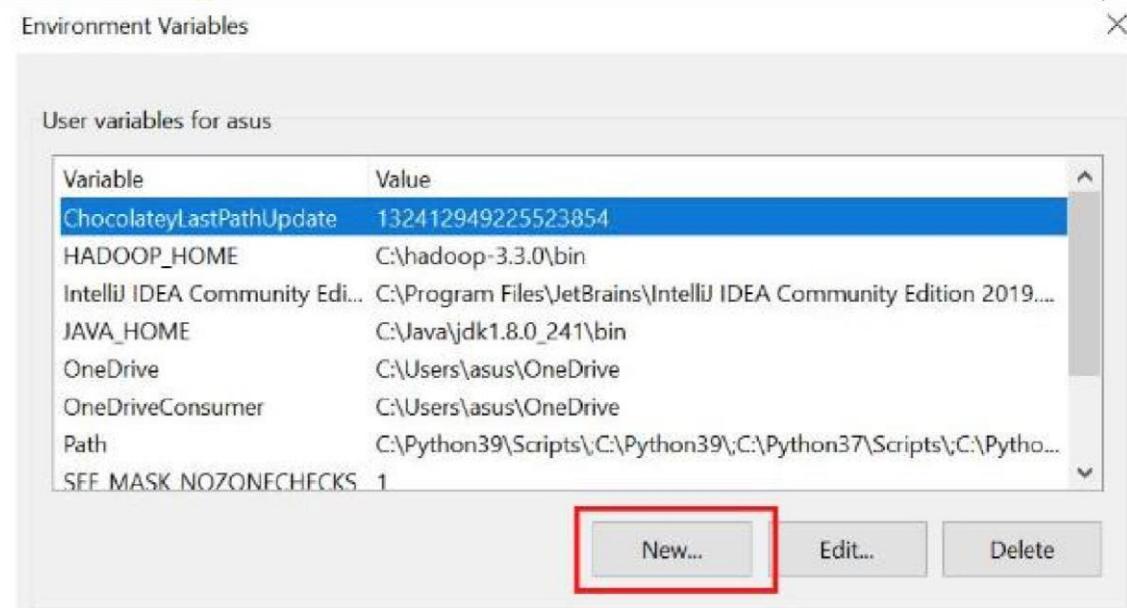
3. Set the path JAVA\_HOME Environment variable
4. Set the path HADOOP\_HOME Environment variable



ssfdgfdgdf

gggggggg





## Environment Variables

| User variables for asus        |   |
|--------------------------------|---|
| Variable                       | Value   |
| ChocolateyLastPathUpdate       | 132412949225523854  |
| HADOOP_HOME                    | C:\hadoop-3.3.0\bin   |
| IntelliJ IDEA Community Edi... | C:\Program Files\JetBrains\IntelliJ IDEA Community Edition 2019.... |
| JAVA_HOME                      | C:\Java\jdk1.8.0_241\bin  |
| OneDrive                       | C:\Users\asus\OneDrive  |
| OneDriveConsumer               | C:\Users\asus\OneDrive  |
| Path                           | C:\Python39\Scripts;C:\Python39;C:\Python37\Scripts;C:\Python...    |
| SFF_MASK_NOZONECHECKS          | 1   |

**New...** **Edit...** **Delete**

## Environment Variables

User variables for asus

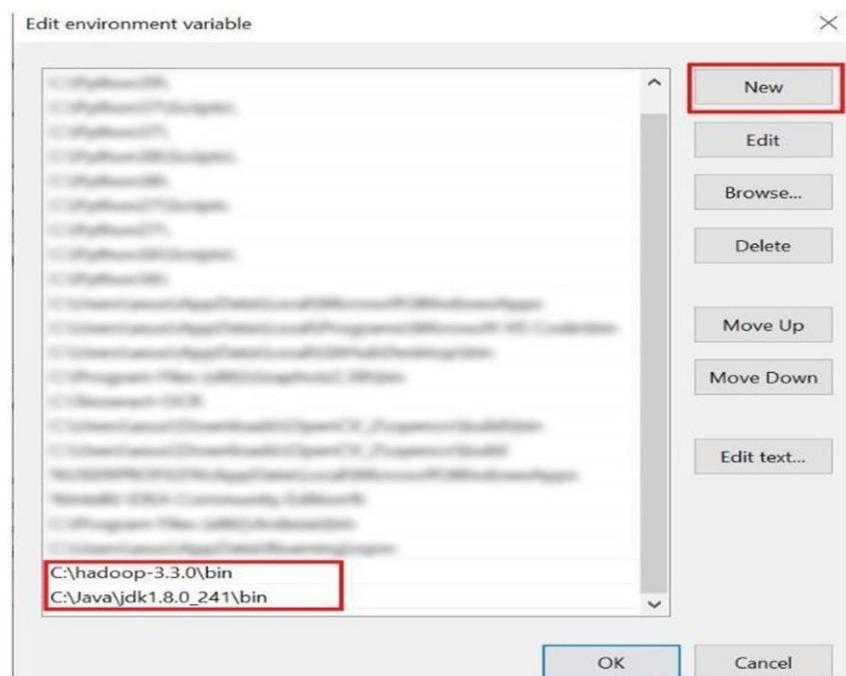
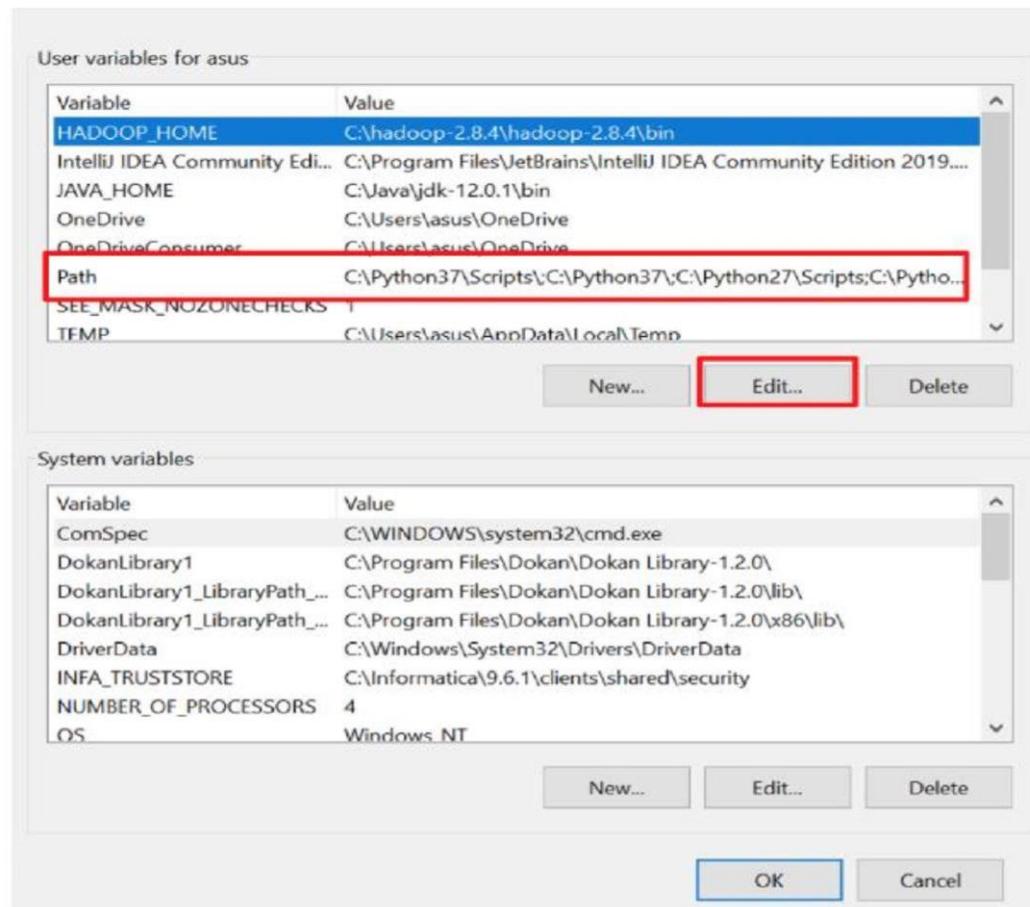
| Variable                       | Value   |
|--------------------------------|---|
| ChocolateyLastPathUpdate       | 132412949225523854  |
| HADOOP_HOME                    | C:\hadoop-3.3.0\bin   |
| IntelliJ IDEA Community Edi... | C:\Program Files\JetBrains\IntelliJ IDEA Community Edition 2019.... |
| JAVA_HOME                      | C:\Java\jdk1.8.0_241\bin  |
| OneDrive                       | C:\Users\asus\OneDrive  |
| OneDriveConsumer               | C:\Users\asus\OneDrive  |
| Path                           | C:\Python39\Scripts;C:\Python39;C:\Python37\Scripts;C:\Python...    |

**Edit User Variable**

Variable name:

Variable value:

**OK** **Cancel**



## 5.Configuration

Edit file C:/Hadoop-3.3.0/etc/hadoop/core-site.xml, paste the xml code in folder and save

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

---

Rename “mapred-site.xml.template” to “mapred-site.xml” and edit this file C:/Hadoop-

3.3.0/etc/hadoop/mapred-site.xml, paste xml code and save this file.

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

---

Create folder “data” under “C:\Hadoop-3.3.0”

Create folder “datanode” under “C:\Hadoop-3.3.0\data”

Create folder “namenode” under “C:\Hadoop-3.3.0\data”

---

Edit file C:\Hadoop-3.3.0/etc/hadoop/hdfs-site.xml, paste xml code and save this file.

```
<configuration>
  <property>
    <name>dfs.replication</name>
```

```
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/hadoop-3.3.0/data/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/hadoop-3.3.0/data/datanode</value>
</property>
</configuration>
```

---

Edit file C:/Hadoop-3.3.0/etc/hadoop/yarn-site.xml, paste xml code and save this file.

```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>
```

---

Edit file C:/Hadoop-3.3.0/etc/hadoop/hadoop-env.cmd by closing the command line

“JAVA\_HOME=%JAVA\_HOME%” instead of set “JAVA\_HOME=C:\Java”

---

---

## 6. Hadoop Configurations

### Download

[https://github.com/brainmentorspvtltd/BigData\\_RDE/blob/master/Hadoop%20Configuration.zip](https://github.com/brainmentorspvtltd/BigData_RDE/blob/master/Hadoop%20Configuration.zip) or (for hadoop 3) <https://github.com/s911415/apache-hadoop-3.1.0-winutils>

- Copy folder bin and replace existing bin folder in C:\Hadoop-3.3.0\bin
- Format the NameNode
- Open cmd and type command “hdfs namenode –format”

```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19041.572]
(c) 2020 Microsoft Corporation. All rights reserved.

C:\hadoop-3.3.0\bin>hdfs namenode -format
```

## 7. Testing

- Open cmd and change directory to C:\Hadoop-3.3.0\sbin □ type start-all.cmd

```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19041.572]
(c) 2020 Microsoft Corporation. All rights reserved.

C:\hadoop-3.3.0\sbin>start-all.cmd
```

(Or you can start like this)

**Start namenode and datanode with this command**

- type start-dfs.cmd

- Start yarn through this command
- type start-yarn.cmd

### Make sure these apps are running

- Hadoop Namenode
- Hadoop datanode
- YARN Resource Manager

#### - YARN Node Manager

```

Apache Hadoop Distribution - hadoop_namenode
Apache Hadoop Distribution - hadoop_datanode
Apache Hadoop Distribution - yarn_resourcemanager
Apache Hadoop Distribution - yarn_nodemanager
Nov 28, 2018 12:21:09 AM com.sun.jersey.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope "Singleton"
Nov 28, 2018 12:21:09 AM com.sun.jersey.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.NMWebServices to GuiceManagedComponentProvider with the scope "Singleton"
The scope "Singleton"
Nov 28, 2018 12:21:09 INFO mortbay.log: Started SelectChannelConnector@0.0.0.0:8042
Nov 28, 2018 12:21:09 INFO webapp.WebApps: Web app node started at 8042
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : DESKTOP-475FCII:61797
Nov 28, 2018 12:21:09 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8090
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: Sending out 0 NM container statuses: []
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: Registering with RM using containers: []
Nov 28, 2018 12:21:09 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id 20336
Nov 28, 2018 12:21:09 INFO security.NMContainerTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id 20336
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as DESKTOP-475FCII:61797 with
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: total resource of <memory:8192, vCores:8>
Nov 28, 2018 12:21:09 INFO nodemanager.NodeStatusUpdaterImpl: Notifying ContainerManager to unblock new container-requests
Nov 28, 2018 12:21:13 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:16 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:19 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:22 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:25 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:28 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:31 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7
Nov 28, 2018 12:21:34 WARN util.SysInfoWindows: Expected split length of sysInfo to be 11. Got 7

```

Open: <http://localhost:8088>

| ID                         | User | Name | Application Type | Queue | Priority | StartTime | FinishTime | Status | FinalStatus | Running Containers | Allocated CPU | Allocated Memory MB | % of Queue | % of Cluster | Progress | Tracking UI | Blacklisted Nodes |
|----------------------------|------|------|------------------|-------|----------|-----------|------------|--------|-------------|--------------------|---------------|---------------------|------------|--------------|----------|-------------|-------------------|
| No data available in table |      |      |                  |       |          |           |            |        |             |                    |               |                     |            |              |          |             |                   |

Showing 0 to 0 of 0 entries

First Previous Next Last

Open: <http://localhost:9870>

localhost:9870/dfshealth.html#tab-overview

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

## Overview 'localhost:9000' (✓active)

|                |  |
|----------------|--|
| Started:       | Sun Nov 08 16:53:46 +0530 2020                             |
| Version:       | 3.3.0 [redacted] 9af                                       |
| Compiled:      | Tue Jul 07 00:14:00 +0530 2020 by brahma from branch-3.3.0 |
| Cluster ID:    | C [redacted]   |
| Block Pool ID: | B [redacted] 44  |

## Summary

Hadoop installed Successfully.....

|                   |  |
|-------------------|--|
| <b>EX. No : 2</b> |  |
| <b>DATE:</b>      | <b>HADOOP IMPLEMENTATION OF FILE MANAGEMENT TASK</b> |

### **AIM:**

To Implementation of hadoop file Management tasks, such as adding files and directories, retrieving files and deleting files

#### **1. Create a directory in HDFS at given path(s).**

Usage:

hadoop fs -mkdir <paths> Example:

hadoop fs -mkdir /user/saurzcode/dir1 /user/saurzcode/dir2

#### **2. List the contents of a directory.**

Usage :

hadoop fs -

ls <args>

Example:

hadoop fs -ls /user/saurzcode

#### **3. Upload and download a file in HDFS.**

**Upload:** hadoop fs -put:

Copy single src file, or multiple src files from local file system to the Hadoop data file system Usage:

hadoop fs -put <localsrc> ... <HDFS\_dest\_Path> Example:

hadoop fs -put /home/saurzcode/Samplefile.txt /user/

saurzcode/dir3/ **Download:** hadoop fs -get:

---

Copies/Downloads files to the local file

system Usage:

hadoop fs -get <hdfs\_src> <localdst> Example:

hadoop fs -get /user/saurzcode/dir3/Samplefile.txt /home/

#### **4. See contents of a file** Same as unix cat command:

Usage:

hadoop fs -cat <path[filename]>

**Example:**

hadoop fs -cat /user/saurzcode/dir1/abc.txt

#### **5. Copy a file from source to destination**

This command allows multiple sources as well in which case the destination must be a directory.

Usage:

hadoop fs -cp <source> <dest>

**Example:**

hadoop        fs        -cp

/user/saurzcode/dir1/abc.txt

/user/saurzcode/ dir2

#### **6. Copy a file from/To Local file system to HDFS copyFromLocal** Usage:

hadoop fs -copyFromLocal <localsrc>

URI Example:

hadoop fs -copyFromLocal /home/saurzcode/abc.txt /user/  
saurzcode/abc.txt Similar to put command, except that the source is  
restricted to a local file reference.

---

**copyToLocal** Usage: hadoop fs -copyToLocal  
[-ignorecrc] [-crc] URI <localdst>

Similar to get command, except that the destination is restricted to a local file reference.

## 7. Move file from source to destination.

Note:- Moving files across filesystem is not permitted.

Usage :

hadoop fs -mv <src> <dest> Example:

hadoop fs -mv /user/saurzcode/dir1/abc.txt /user/saurzcode/ dir2

## 8. Remove a file or directory in HDFS.

Remove files specified as argument. Deletes directory only when it is empty Usage :

hadoop fs -rm <arg> Example:

hadoop fs -rm /user/saurzcode/dir1/abc.txt

### Recursive version of delete.

Usage :

hadoop fs -rmr <arg>

Example: hadoop fs -rmr  
/user/saurzcode/

## 9. Display last few lines of a file.

Similar to tail command in Unix.

Usage :

---

```
hadoop fs -tail <path[filename]>
```

Example: hadoop fs -tail

```
/user/saurzcode/dir1/abc.txt
```

#### **10. Display the aggregate length of a file.**

Usage :

```
hadoop fs -du <path> Example:
```

```
hadoop fs -du /user/saurzcode/dir1/abc.txt
```

#### **RESULT:**

Thus the Implementation of hadoop file Management tasks, such as adding files and directories, retrieving files and deleting files has been executed successfully.

|                   |   |
|-------------------|---|
| <b>EX. No : 3</b> | <b>IMPLEMENT OF MATRIX MULTIPLICATION WITH HADOOP</b> |
| <b>DATE:</b>      | <b>MAP REDUCE</b>                                     |

### **AIM:-**

To write a Map Reduce Program that implements Matrix Multiplication.

### **PROCEDURE:**

We assume that the input matrices are already stored in Hadoop Distributed File System (HDFS) in a suitable format (e.g., CSV, TSV) where each row represents a matrix element. The matrices are compatible for multiplication (the number of columns in the first matrix is equal to the number of rows in the second matrix).

#### **STEP 1: MAPPER**

The mapper will take the input matrices and emit key-value pairs for each element in the result matrix. The key will be the (row, column) index of the result element, and the value will be the corresponding element value.

#### **STEP 2: REDUCER**

The reducer will take the key-value pairs emitted by the mapper and calculate the partial sum for each element in the result matrix.

#### **STEP 3: MAIN DRIVER**

The main driver class sets up the Hadoop job configuration and specifies the input and output paths for the matrices.

#### **STEP 4: RUNNING THE JOB**

To run the MapReduce job, you need to package your classes into a JAR file and then submit it to Hadoop using the hadoop jar command. Make sure to replace `input_path` and `output_path` with the actual HDFS paths to your input matrices and desired output directory.

---

## **PROGRAM:**

```
import java.io.IOException; import
java.util.StringTokenizer; import
org.apache.hadoop.io.IntWritable;
import
org.apache.hadoop.io.LongWritabl
e; import
org.apache.hadoop.io.Text; import
org.apache.hadoop.mapreduce.Ma
pper; import
org.apache.hadoop.mapreduce.Red
ucer; import
org.apache.hadoop.conf.Configurat
ion; import
org.apache.hadoop.mapreduce.Job;
import
org.apache.hadoop.mapreduce.lib.i
nput.TextInputFormat; import
org.apache.hadoop.mapreduce.lib.
output.TextOutputFormat; import
org.apache.hadoop.mapreduce.lib.i
nput.FileInputFormat; import
org.apache.hadoop.mapreduce.lib.
output.FileOutputFormat; import
org.apache.hadoop.fs.Path; public
class MatrixMultiplicationMapper
extends Mapper<LongWritable,
Text, Text, Text>

{

protected void map(LongWritable key, Text value, Context context) throws
IOException,
```

```

InterruptedException {

    // Parse the input line to get row, column, and value of each element in the
    // input matrices      String[] elements = value.toString().split(",");
    //           int row =
    Integer.parseInt(elements[0]);
    //           int col = Integer.parseInt(elements[1]);
    //           int
    val = Integer.parseInt(elements[2]);

    // Emit key-value pairs where key is (row, column) index of the result element
    // and value is the corresponding element value
    context.write(new Text(row + "," + col), new Text(val));

}

} public class MatrixMultiplicationReducer extends Reducer<Text, Text, Text,
IntWritable> { protected void reduce(Text key, Iterable<Text> values, Context
context) throws IOException, InterruptedException {      int result = 0;
for
(Text value : values) {

    // Accumulate the partial sum for the result
    element      result +=
    Integer.parseInt(value.toString());
}

// Emit the final result for the result element
context.write(key, new IntWritable(result));

}
} public class
MatrixMultiplicationDriver {

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Matrix Multiplication");
    job.setJarByClass(MatrixMultiplicationDriver.class);
    job.setMapperClass(MatrixMultiplicationMapper.class);
}

```

```
job.setReducerClass(MatrixMultiplicationReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(Text.class);

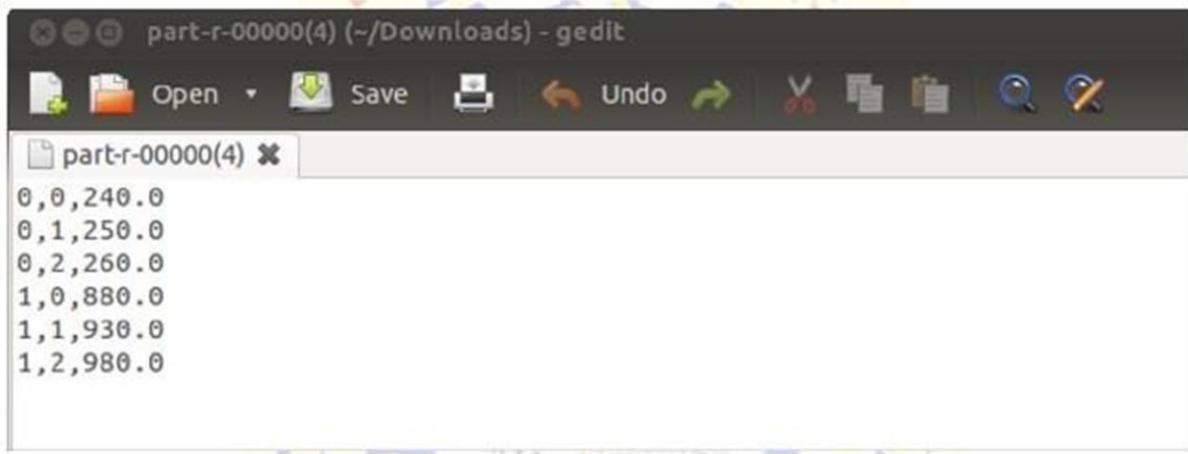
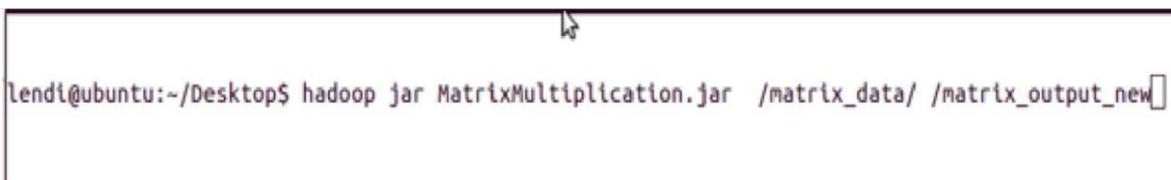
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}
```

**Run the program** hadoop jar matrixmultiplication.jar

MatrixMultiplicationDriver input\_path output\_path

```
lendi@ubuntu:~/Desktop$ hadoop jar MatrixMultiplication.jar /matrix_data/ /matrix_output_new
```



The screenshot shows a terminal window with the command `lendi@ubuntu:~/Desktop$ hadoop jar MatrixMultiplication.jar /matrix_data/ /matrix_output_new` entered and executed. Below the terminal is a gedit text editor window titled "part-r-00000(4) (~/Downloads) - gedit". The file contains the following text:

```
0,0,240.0
0,1,250.0
0,2,260.0
1,0,880.0
1,1,930.0
1,2,980.0
```

---

## **RESULT:**

Thus the Map Reduce Program that implements Matrix Multiplication was executed and verified successfully.

|                  |  |
|------------------|--|
| <b>EX. NO: 4</b> | <b>RUN A BASIC WORD COUNT MAP REDUCE PROGRAM</b> |
| <b>DATE:</b>     |  |

**AIM:-**

To write a Basic Word Count program to understand Map Reduce Paradigm.

**PROCEDURE:**

The entire MapReduce program can be fundamentally divided into three parts:

- Mapper Phase Code
- Reducer Phase Code
- Driver Code

**STEP 1: MAPPER CODE:**

We have created a class Map that extends the class Mapper which is already defined in the MapReduce Framework.

- We define the data types of input and output key/value pair after the class declaration using angle brackets.
- Both the input and output of the Mapper is a key/value pair.

**Input:**

- The key is nothing but the offset of each line in the text file:LongWritable  
The value is each individual line : Text

- The key is the tokenized words: Text
- We have the hardcoded value in our case which is 1: IntWritable
- **Example** – Dear 1, Bear 1, etc.

We have written a java code where we have tokenized each word and assigned them a hardcoded value equal to 1.

**STEP 2 : REDUCER CODE:**

- We have created a class Reduce which extends class Reducer like that of Mapper.

- We define the data types of input and output key/value pair after the class declaration using angle brackets as done for Mapper.
- Both the input and the output of the Reducer is a key value pair.

#### **Input:**

- The key nothing but those unique words which have been generated after the sorting and shuffling phase: Text
- The value is a list of integers corresponding to each key: IntWritable □ Example – Bear, [1, 1], etc.

#### **Output:**

- The key is all the unique words present in the input text file: Text
- The value is the number of occurrences of each of the unique words: IntWritable □ Example – Bear, 2; Car, 3, etc.
- We have aggregated the values present in each of the list corresponding to each key and produced the final answer.
- In general, a single reducer is created for each of the unique words, but, you can specify the number of reducer in mapred-site.xml.

#### **STEP 3: DRIVER CODE:**

- In the driver class, we set the configuration of our MapReduce job to run in Hadoop.
- We specify the name of the job , the data type of input/ output of the mapper and reducer.
- We also specify the names of the mapper and reducer classes.
- The path of the input and output folder is also specified.
- The method setInputFormatClass () is used for specifying that how a Mapper will read the input data or what will be the unit of work. Here, we have chosen TextInputFormat so that single line is read by the mapper at a time from the

---

input text file. The main () method is the entry point for the driver. In this method, we instantiate a new Configuration object for the job.

### **PROGRAM:**

```
import java.io.IOException; import  
java.util.StringTokenizer; import  
org.apache.hadoop.io.IntWritable;  
import  
org.apache.hadoop.io.LongWritabl  
e; import  
org.apache.hadoop.io.Text; import  
org.apache.hadoop.mapreduce.Map  
per; import  
org.apache.hadoop.mapreduce.Red  
ucer; import  
org.apache.hadoop.conf.Configurat  
ion; import  
org.apache.hadoop.mapreduce.Job;  
import  
org.apache.hadoop.mapreduce.lib.i  
nput.TextInputFormat; import  
org.apache.hadoop.mapreduce.lib.o  
utput.TextOutputFormat; import  
org.apache.hadoop.mapreduce.lib.i  
nput.FileInputFormat; import  
org.apache.hadoop.mapreduce.lib.o  
utput.FileOutputFormat; import  
org.apache.hadoop.fs.Path; public  
class WordCount  
  
{ public static class Map extends  
Mapper<LongWritable,Text,Text,IntWritable> { public void  
map(LongWritable key, Text value,Context context) throws
```

```
    } catch (IOException | InterruptedException e) {
        e.printStackTrace();
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "My Word Count Program");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(WordCountMapper.class);
    job.setReducerClass(WordCountReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

```
job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);

Path outputPath = new Path(args[1]);
//Configuring the input/output path from the filesystem into the job
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
//deleting the output path automatically from hdfs so that we
don't have to delete it explicitly
outputPath.getFileSystem(conf).delete(outputPath); //exiting
the job only if the flag value becomes false

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

**Run the MapReduce code:**

**The command for running a MapReduce code is:**

```
hadoop jar hadoop-mapreduce-example.jar WordCount /sample/input
/samples/output
```

```
lendi@ubuntu: ~/Desktop
16/08/17 01:17:45 INFO impl.YarnClientImpl: Submitted application application_1471410736896_0001
16/08/17 01:17:45 INFO mapreduce.Job: The url to track the job: http://ubuntu.ubuntu-domain:8088/proxy/application_1471410736896_0001/
16/08/17 01:17:45 INFO mapreduce.Job: Running job: job_1471410736896_0001
16/08/17 01:17:52 INFO mapreduce.Job: Job job_1471410736896_0001 running in uber mode : false
16/08/17 01:17:52 INFO mapreduce.Job: map 0% reduce 0%
16/08/17 01:17:59 INFO mapreduce.Job: map 100% reduce 0%
16/08/17 01:18:06 INFO mapreduce.Job: map 100% reduce 100%
16/08/17 01:18:06 INFO mapreduce.Job: Job job_1471410736896_0001 completed successfully
16/08/17 01:18:06 INFO mapreduce.Job: Counters: 49
File System Counters
    FILE: Number of bytes read=3772644
    FILE: Number of bytes written=7775215
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=174718
    HDFS: Number of bytes written=510970
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
```

part-r-00000(3) (~/Downloads) - gedit

| Word         | Count |
|--------------|-------|
| 2.           | 1     |
| 3.           | 28    |
| 3.           | 1     |
| 4.           | 1     |
| 5.           | 1     |
| 6.           | 1     |
| 7.           | 1     |
| 8.           | 1     |
| 9.           | 1     |
| A            | 1012  |
| A'           | 2     |
| ADRIAN,      | 2     |
| AEdiles,     | 1     |
| AEsculapius? | 1     |
| ALARBUS,     | 1     |
| ALENCON,     | 2     |
| ALL'S        | 25    |
| ANDRONICUS,  | 1     |
| ANGELO,      | 2     |

## RESULT:

Thus the Map Reduce Program that implements word count was executed and verified successfully.

|            |  |
|------------|--|
| EX. NO : 5 | INSTALLATION OF HIVE ALONG WITH PRACTICE |
| DATE:      |  |

### PREREQUISITES:

- Java Development Kit (JDK) installed and the JAVA\_HOME environment variable set.
- Hadoop installed and configured on your Windows system.

### STEP-BY-STEP INSTALLATION:

#### 1. Download HIVE:

Visit the Apache Hive website and download the latest stable version of Hive.

Official Apache Hive website: <https://hive.apache.org/>

#### 2. Extract the Downloaded Hive Archive to a Directory on Your Windows Machine, e.g., C:\hive.

#### 3. Configure Hive:

- Open the Hive configuration file (hive-site.xml) located in the conf folder of the extracted Hive directory.
- Set the necessary configurations, such as Hive Metastore connection settings and Hadoop configurations. Make sure to adjust paths accordingly for Windows. Here's an example of some configurations:

```
<configuration>
<property>
  <name>javax.jdo.option.ConnectionURL</name>
  <value>jdbc:derby:/path/to/metastore_db;create=true</value>
```

```
<description>JDBC connect string for a JDBC metastore.</description>
</property>

<!-- Other Hive configurations -->
</configuration>
```

#### 4. Environment Variables Setup:

- Add the Hive binary directory (C:\hive\bin in this example) to your PATH environment variable.
- Set the HIVE\_HOME environment variable to point to the Hive installation directory (C:\hive in this example).

#### 5. Start the Hive Metastore service:

To start the Hive Metastore service, you can use the schematool script:

```
bash
Copy code

schematool -initSchema -dbType derby
```

#### 6. Start Hive:

- Open a command prompt or terminal and navigate to the Hive installation directory.
- Execute the hive command to start the Hive shell.

#### EXAMPLES:

##### 1. Create a Database:

To create a new database in HIVE, use the following syntax:

```
CREATE DATABASE
database_name; Example:

CREATE DATABASE mydatabase;
```

##### 2. Use a Database:

To use a specific database in HIVE, use the following syntax: USE database\_name; Example:

---

```
USE mydatabase;
```

### 3. Show Databases:

To display a list of available databases in HIVE, use the following syntax:

```
SHOW DATABASES;
```

### 4. Create a Table:

To create a table in HIVE, use the following

```
syntax: CREATE TABLE table_name (
```

```
    column1
```

```
    datatype,
```

```
    column2
```

```
    datatype,
```

```
    ...
```

```
);
```

#### Example:

```
CREATE TABLE mytable (
```

```
    id INT,
```

```
    name STRING,
```

```
    age INT
```

```
);
```

### 5. Show Tables:

To display a list of tables in the current database, use the following syntax:

```
SHOW TABLES;
```

### 6. Describe a Table:

To view the schema and details of a specific table, use the following

```
syntax: DESCRIBE table_name;
```

**Example:**

```
DESCRIBE mytable;
```

## 7. Insert Data into a Table:

To insert data into a table in HIVE, use the following syntax:

```
INSERT INTO table_name (column1, column2, ...) VALUES (value1, value2,  
...);
```

**Example:**

```
INSERT INTO mytable (id, name, age) VALUES (1, 'John Doe', 25);
```

## 8. Select Data from a Table:

```
SELECT column1, column2, ... FROM table_name WHERE  
condition;
```

**Example:**

```
SELECT * FROM mytable WHERE age > 20;
```

## RESULT:

Thus the Installation of HIVE was done successfully.

|                   |   |
|-------------------|---|
| <b>EX. NO : 6</b> | <b>INSTALLATION OF HBASE ALONG WITH PRACTICE EXAMPLES</b> |
| <b>DATE:</b>      |   |

### **AIM:**

To install HBASE using Virtual Machine and perform some operations in HBASE.

### **PROCEDURE:**

#### **Step 1: Install a Virtual Machine**

- Download and install a virtual machine software such as VirtualBox (<https://www.virtualbox.org/>) or VMware (<https://www.vmware.com/>).
- Create a new virtual machine and install a Unix-based operating system like Ubuntu or CentOS. You can download the ISO image of your desired Linux distribution from their official websites.

#### **Step 2: Set up the Virtual Machine**

- Launch the virtual machine and install the Unix-based operating system following the installation wizard.
- Make sure the virtual machine has network connectivity to download software packages.

#### **Step 3: Install Java**

- Open the terminal or command line in the virtual machine.
- Update the package list **sudo apt update**
- Install OpenJDK (Java Development Kit):  
**sudo apt install default-jdk**

Verify the Java

installation: **java -version**

#### **Step 4: Download and Install HBase**

- In the virtual machine, navigate to the directory where you want to install HBase.
- Download the HBase binary distribution from the Apache HBase website (<https://hbase.apache.org/>). Look for the latest stable version.
- Extract the downloaded archive **tar -xvf <hbase\_archive\_name>.tar.gz**
- Replace `<hbase_archive_name>` with the actual name of the HBase archive file.
- Move the extracted HBase directory to a desired location:  
**sudo mv <hbase\_extracted\_directory> /opt/hbase**
- Replace `<hbase_extracted_directory>` with the actual name of the extracted HBase directory.

#### **Step 5: Configure HBase**

- Open the HBase configuration file for editing: **sudo nano /opt/hbase/conf/hbase-site.xml**
- Add the following properties to the configuration file:

```
<configuration>
  <property>
    <name>hbase.rootdir</name>
    <value>file:///var/lib/hbase</value>
  </property>
  <property>
    <name>hbase.zookeeper.property.dataDir</name>
    <value>/var/lib/zookeeper</value>
  </property>
</configuration>
```

- Save the file and exit the text editor.

#### **Step 6: Start HBase**

- Start the HBase server: **sudo /opt/hbase/bin/start-hbase.sh**

---

## **HBASE PRACTICE EXAMPLES:**

### **Step 1: Start HBase**

- Make sure HBase is installed and running on your Windows system.

### **Step 2: Open HBase Shell**

- Open a command prompt or terminal window and navigate to the directory where the HBase installation is located. Run the following command to start the HBase shell:

```
>>hbase shell
```

### **Step 3: Create a Table**

- In the HBase shell, you can create a table with column families.
- For example, let's create a table named "my\_table" with a column family called "cf":

```
>> create 'my_table', 'cf'
```

### **Step 4: Insert Data**

- To insert data into the table, you can use the put command.
- Here's an example of inserting a row with a specific row key and values:

```
>> put 'my_table', 'row1', 'cf:column1', 'value1'  
>> put 'my_table', 'row1', 'cf:column2', 'value2'
```

### **Step 5: Get Data**

- You can retrieve data from the table using the get command.
- For example, to get the values of a specific row:

```
>> get 'my_table', 'row1'
```

- This will display all the column family values for the specified row.

### **Step 6: Scan Data**

- To scan and retrieve multiple rows or the entire table, use the scan command.
- For instance, to scan all rows in the table:

```
>> scan 'my_table'
```

- This will display all rows and their corresponding column family values.

### **Step 7: Delete Data**

- 
- To delete a specific row or a particular cell value, you can use the delete command.
  - Here's an example of deleting a specific row:

```
>>delete 'my_table', 'row1'
```

### **Step 8: Disable and Drop Table**

- If you want to remove the table entirely, you need to disable and drop it.
- Use the following commands:

```
>>disable 'my_table'
```

```
>>drop 'my_table'
```

### **RESULT:**

Thus the installation of HBase using Virtual Machine was done successfully.

|                   |                               |
|-------------------|-------------------------------|
| <b>EX. NO : 7</b> | <b>INSTALLATION OF THRIFT</b> |
| <b>DATE:</b>      |                               |

## **AIM:**

To install Apache thrift on Windows OS.

## **PROCEDURE:**

### **Step 1: Download Apache Thrift:**

- Visit the Apache Thrift website: <https://thrift.apache.org/>
- Go to the "Downloads" section and find the latest version of Thrift.
- Download the Windows binary distribution (ZIP file) for the desired version.

### **Step 2: Extract the ZIP file:**

- Locate the downloaded ZIP file and extract its contents to a directory of your choice.
- This directory will be referred to as <THRIFT\_DIR> in the following steps.

### **Step 3: Set up environment variables:**

- Open the Start menu and search for "Environment Variables" and select "Edit the system environment variables."
- Click the "Environment Variables" button at the bottom right of the "System Properties" window.
- Under the "System variables" section, find the "Path" variable and click "Edit."
- Add the following entries to the "Variable value" field (replace <THRIFT\_DIR> with the actual directory path):

**<THRIFT\_DIR>\bin**

---

<THRIF

**T\_DIR>\lib** □ Click "OK"

to save the changes.

**Step 4: Verify the installation:**

- Open a new Command Prompt window.
- Run the following command to verify that Thrift is installed and accessible:  
**thrift –version**
- If everything is set up correctly, you should see the version number of Thrift printed on the screen.

**RESULT:**

Thus the installation of Thrift on windows OS was done successfully.

|                   |
|-------------------|
| <b>EX. NO : 8</b> |
| <b>DATE:</b>      |

## **PRACTICE IMPORTING AND EXPORTING DATA FROM VARIOUS DATABASES**

### **AIM:**

To import and export data from various Databases using SQOOP.

### **PROCEDURE:**

#### **Step 1: Install SQOOP.**

- First, you need to install Sqoop on your Hadoop cluster or machine.
- Download the latest version of Sqoop from the Apache Sqoop website (<http://sqoop.apache.org/>) and follow the installation instructions provided in the documentation

#### **Step 2: Importing data from a database:**

- To import data from a database into Hadoop, use the following Sqoop command:

```
Sqoop           import          --connect
jdbc:<DB_TYPE>://<DB_HOST>:<DB_PORT>/<DB_NAME>\

--username <DB_USERNAME> \
--password <DB_PASSWORD> \
--table <TABLE_NAME> \
--target-dir <HDFS_TARGET_DIR> \
--m <NUMBER_OF_MAP_TASKS>
```

- Replace the placeholders
- (<DB\_TYPE>, <DB\_HOST>, <DB\_PORT>, <DB\_NAME>, <DB\_USERNAME>,

---

<DB\_PASSWORD>, <TABLE\_NAME>, <HDFS\_TARGET\_DIR>, and <NUMBER\_OF\_MAP\_TASKS>) with the appropriate values for your database and Hadoop environment.

### Step 3: Exporting data to a database:

To export data from Hadoop to a database, use the following Sqoop command:

```
sqoop export --connect  
  jdbc:<DB_TYPE>://<DB_HOST>:<DB_PORT>/<DB_NAME>\  
    --username <DB_USERNAME> \  
    --password <DB_PASSWORD> \  
    --table <TABLE_NAME> \  
    --export-dir <HDFS_EXPORT_DIR> \  
    --input-fields-terminated-by '<DELIMITER>'
```

- Replace the placeholders
- (<DB\_TYPE>, <DB\_HOST>, <DB\_PORT>, <DB\_NAME>, <DB\_USERNAME>, <DB\_PASSWORD>, <TABLE\_NAME>, <HDFS\_EXPORT\_DIR>, and <DELIMITER>) with the appropriate values for your database and Hadoop environment.

---

**RESULT:**

Thus the implementation export data from various Databases using SQOOP was done successfully.