

ANNAI MIRA COLLEGE OF ENGINEERING AND TECHNOLOGY

NH-46, Chennai-Bengaluru National Highways, Arappakkam,

Ranipet-632517, Tamil Nadu, India

Telephone: 04172-292925 Fax: 04172-292926

Email:amcet.rtet@gmail.com/info@amcet.in Web: www.amcet.in

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING



CCS341 – DATA WAREHOUSING LABORATORY

Name :

Register Number :.....

Year & Branch :.....

Semester :.....

Academic Year :.....

ANNAI MIRA COLLEGE OF ENGINEERING AND TECHNOLOGY

NH-46, Chennai-Bengaluru National Highways, Arappakkam,

Ranipet-632517, TamilNadu, India

Telephone: 04172-292925 Fax: 04172-292926



CERTIFICATE

This is to Certify that the Bonafide record of the practical work done by.....
Register Number..... of IIIrd year B.E (Computer Science and Engineering)
submitted for the B.E-Degree practical examination(VIth Semester) in **CCS341 – DATA WAREHOUSING LABORATORY** during the academic year **2023 – 2024**.

Staff in –Charge

Head of the Department

Submitted for the practical examination held on -----

Internal Examiner

External Examiner

TABLE OF CONTENT

SL.NO	DATE	LIST OF EXPERIMENTS	PG.NO	SIGNATURE
1.		Data exploration and integration with WEKA	1	
2.		Apply WEKA tool for data validation	6	
3.		Plan the architecture for real time application	9	
4(A).		Query for star schema using sql server management studio	17	
4(B).		Query for snowflake schema using sql server management studio	20	
5.		Design data warehouse for real time applications	23	
6.		Analyse the Dimensional Modelling	26	
7.		Case study using OLAP	30	
8.		Case study using OLTP	36	
9.		Implementation of warehouse testing	39	

EX NO: 01

DATA EXPLORATION AND INTEGRATION WITH WEKA

DATE:

Aim:

The goal of this lab is to install and familiarize with Weka. To demonstrate the available features in pre-processing, we will use the Weather dataset.

Procedure:

Step1: Download and install Weka.

Step2: Open Weka and have a look at the interface. It is an open-source project written in Java from the University of Waikato.

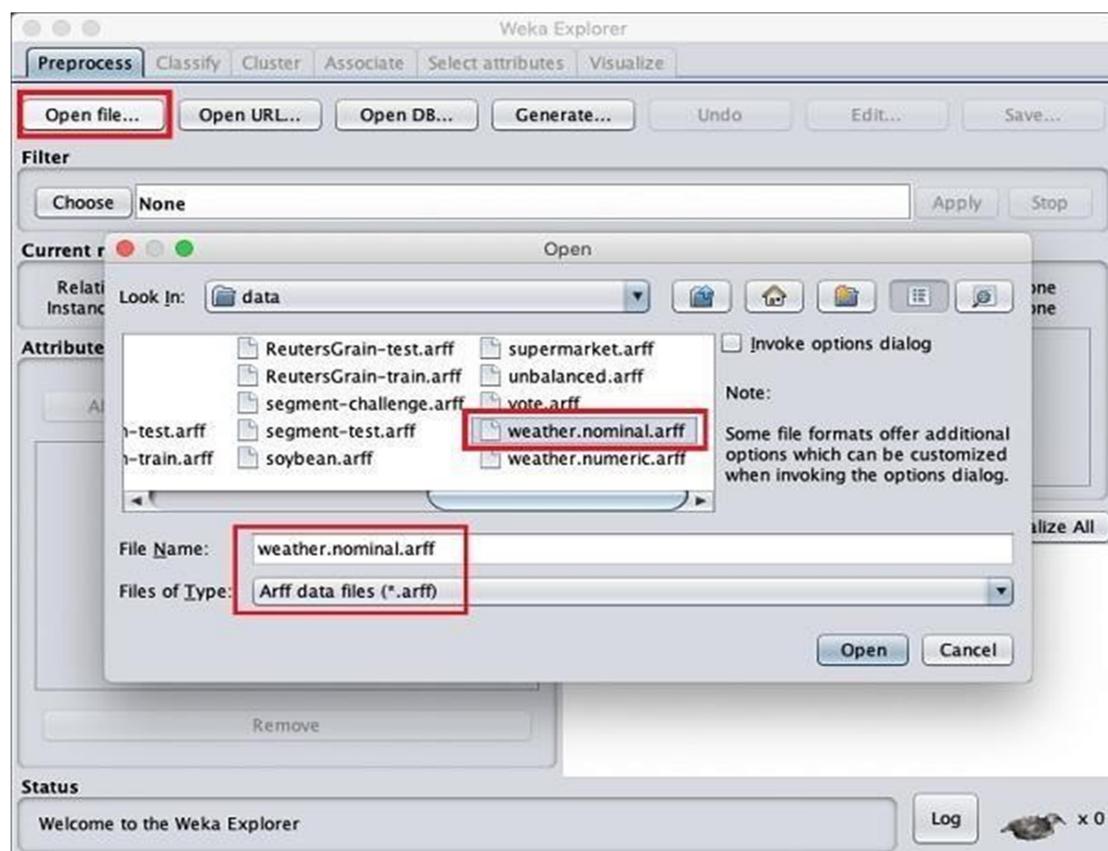
Step 3: Click on the Explorer button on the right side.



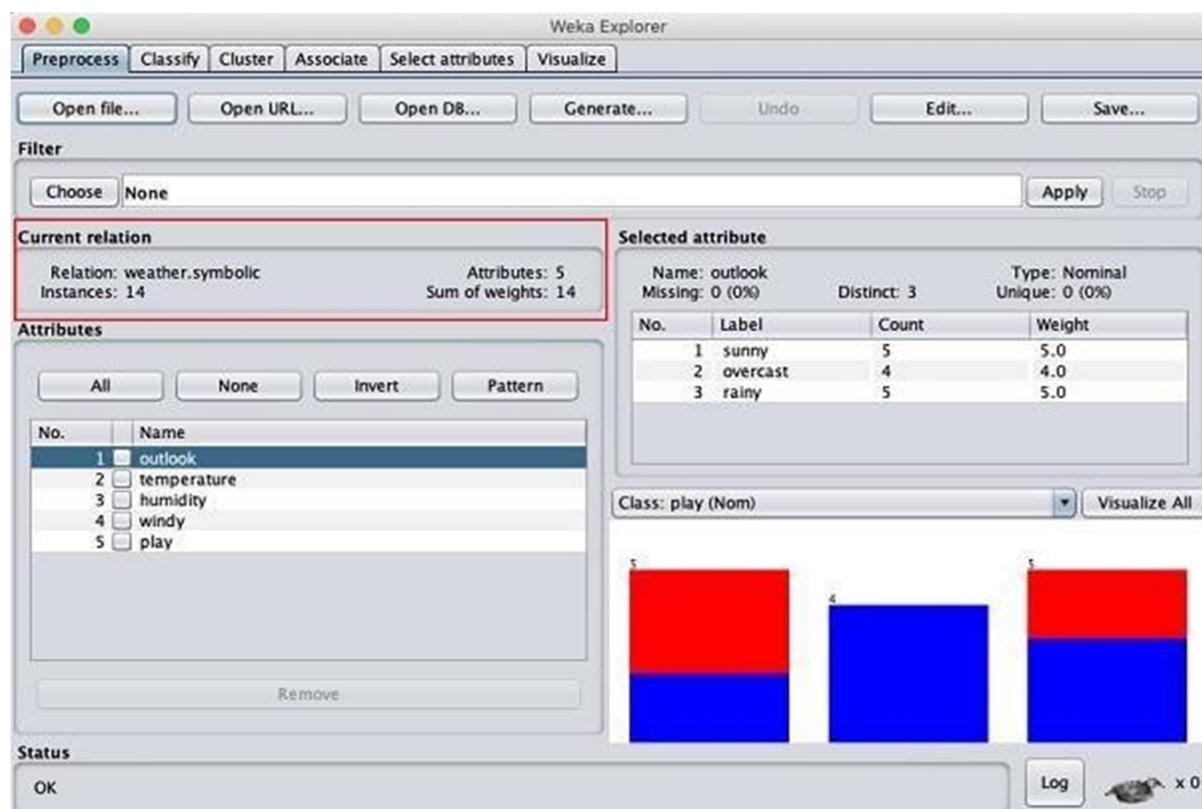
Step 4: Weka comes with a number of small datasets. Those files are located at C:\Program Files\Weka3-8 (If it is installed at this location. Or else, search for Weka-3-8 to find the installation location).

In this folder, there is a subfolder named 'data'. Open that folder to see all files that comes with Weka.

Using the ... **Open file** option under the **Pre-process** tag select the **weather-nominal.arff** file.



When opening the file, the screen looks like this.



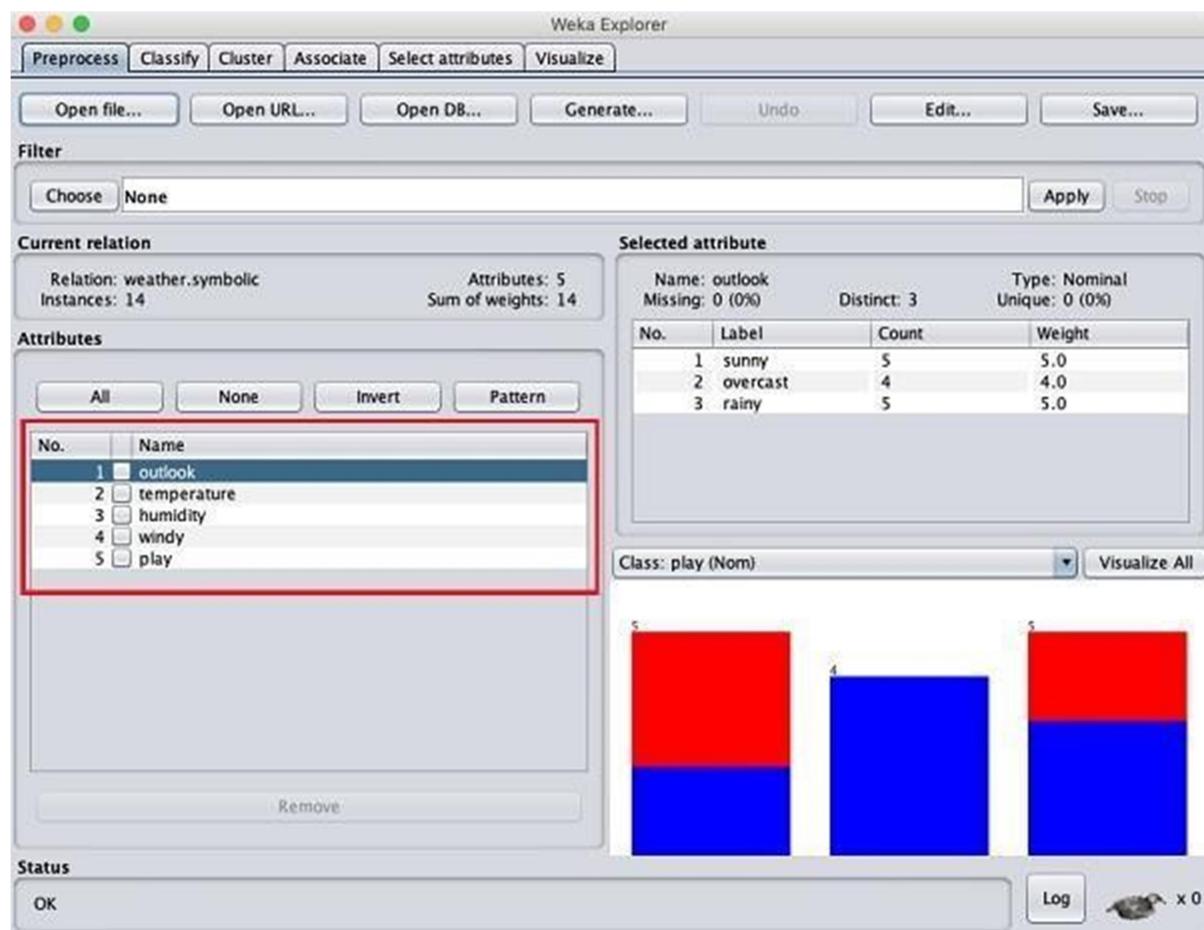
Step 5: Check different tabs to familiarize with the tool.

Understanding Data

Let us first look at the highlighted **Current relation** sub window. It shows the name of the dataset that is currently loaded. You can infer two points from this sub window.

- There are 14 instances – the number of rows in the table.
- The table contains 5 attributes – the fields, which are discussed in the upcoming sections.

On the left side, notice the **Attributes** sub window that displays the various fields in the database.

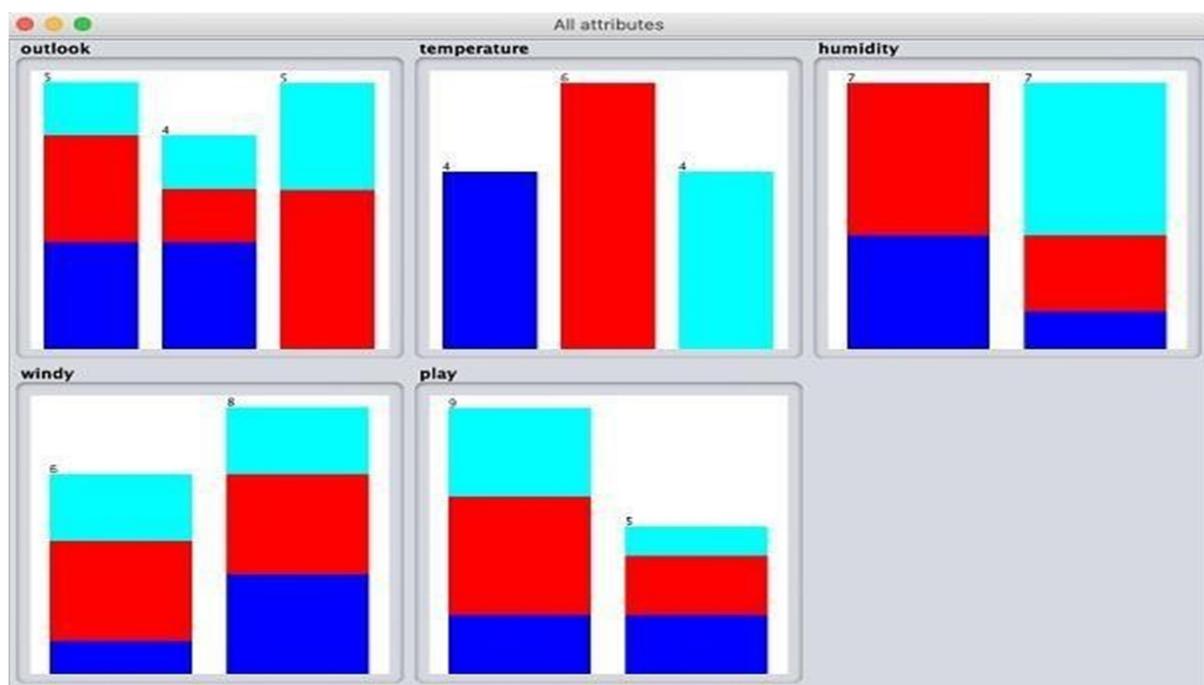


In the **Selected Attribute** sub window, you can observe the following –

- The name and the type of the attribute are displayed.
- The type for the **temperature** attribute is **Nominal**.
- The number of **Missing** values is zero.

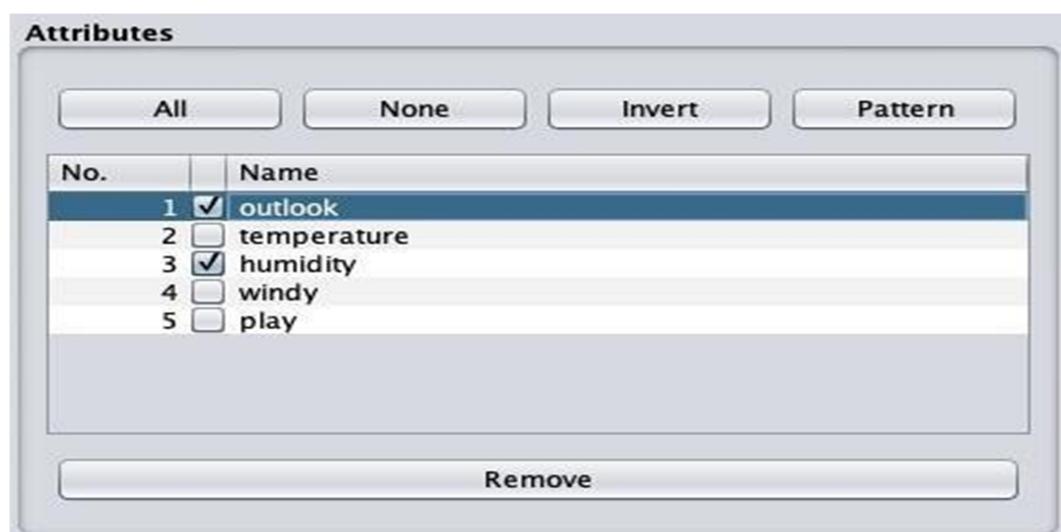
- There are three distinct values with no unique value.
- The table underneath this information shows the nominal values for this field as hot, mild and cold.
- It also shows the count and weight in terms of a percentage for each nominal value.

At the bottom of the window, you see the visual representation of the **class** values. If you click on the **Visualize All** button, you will be able to see all features in one single window as shown here



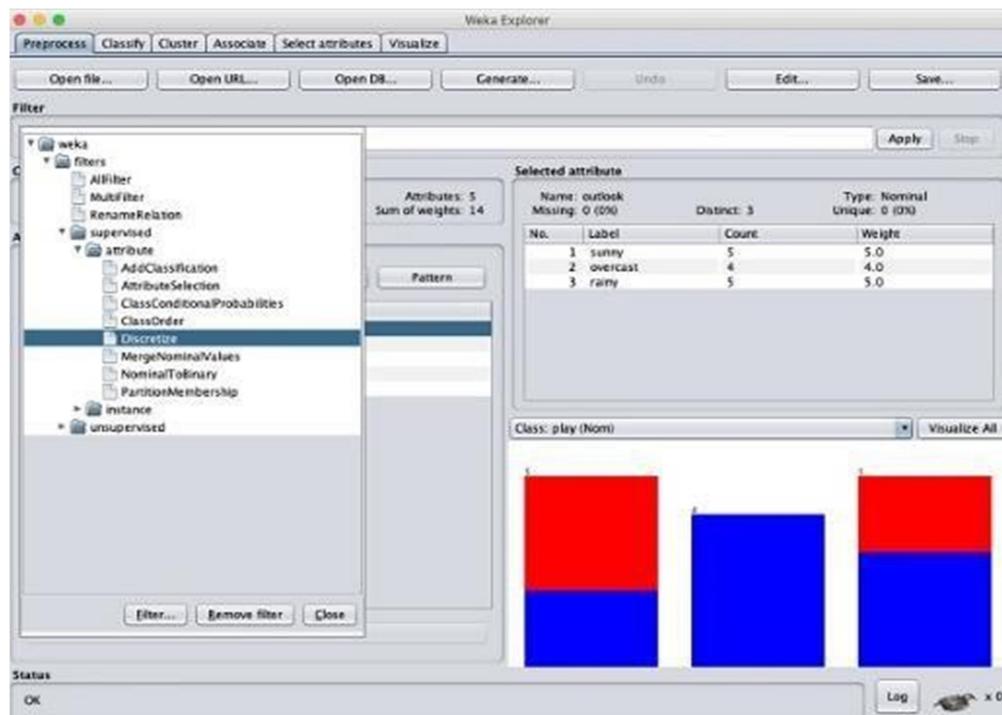
Removing Attributes:

Many a time, the data that you want to use for model building comes with many irrelevant fields. For example, the customer database may contain his mobile number which is relevant in analysing his credit rating.



Applying Filters:

Some of the machine learning techniques such as association rule mining requires categorical data. To illustrate the use of filters, we will use **weather-numeric.arff** database that contains two **numeric** attributes - **temperature** and **humidity**.



Result:

Thus, the data exploration and integration with WEKA is successfully executed.

EX NO: 02

APPLY WEKA TOOL FOR DATA VALIDATION

DATE:

Aim:

To implement data validation using Weka.

Procedure:

Step 1: Launch Weka Explorer

- Open Weka and select the "Explorer" from the Weka GUI Chooser.

Step 2: Load the dataset

- Click on the "Open file" button and select "datasets" > "iris.arff" from the Weka installation directory. This will load the Iris dataset.

Step 3: Split your data into training and testing sets. Under the "Classify" tab, click on the "Choose" button next to the "Test options" area and select a testing method. Weka offers options like cross-validation, percentage split, and user-defined test set. Configure the options according to your needs.

Step 4: Select a classifier algorithm. Weka offers a wide range of algorithms for classification, regression, clustering, and other tasks. Under the "Classify" tab, click on the "Choose" button next to the "Classifier" area and choose an algorithm. Configure its parameters, if needed.

Step 5: Click on the "Start" button under the "Classify" tab to run the training and testing process. Weka will train the model on the training set and test its performance on the testing set using the selected algorithm.

Validation Techniques:

Cross-Validation: Go to the "Classify" tab and choose a classifier. Then, under the "Test options," select the type of cross-validation you want to perform (e.g., 10-fold cross validation). Click "Start" to run the validation.

Train-Test Split: You can also split your data into a training set and a test set. Use the "Supervised" tab to train a model on the training set and evaluate its performance on the test set.

Output

The screenshot shows the Oracle Database SQL Developer interface. A table named 'MATERIAL' is displayed with the following data:

ID	Description	2nd article ID	3rd article ID	4th position	5th class
1	1	2	3	4	01_Fruit
2	2	3	4	5	02_Fruit
3	3	4	5	6	03_Fruit
4	4	5	6	7	04_Fruit
5	5	6	7	8	05_Fruit
6	6	7	8	9	06_Fruit
7	7	8	9	10	07_Fruit
8	8	9	10	11	08_Fruit
9	9	10	11	12	09_Fruit
10	10	11	12	13	10_Fruit
11	11	12	13	14	11_Fruit
12	12	13	14	15	12_Fruit
13	13	14	15	16	13_Fruit
14	14	15	16	17	14_Fruit
15	15	16	17	18	15_Fruit
16	16	17	18	19	16_Fruit
17	17	18	19	20	17_Fruit
18	18	19	20	21	18_Fruit
19	19	20	21	22	19_Fruit
20	20	21	22	23	20_Fruit
21	21	22	23	24	21_Fruit
22	22	23	24	25	22_Fruit
23	23	24	25	26	23_Fruit
24	24	25	26	27	24_Fruit
25	25	26	27	28	25_Fruit
26	26	27	28	29	26_Fruit
27	27	28	29	30	27_Fruit
28	28	29	30	31	28_Fruit
29	29	30	31	32	29_Fruit
30	30	31	32	33	30_Fruit
31	31	32	33	34	31_Fruit
32	32	33	34	35	32_Fruit
33	33	34	35	36	33_Fruit
34	34	35	36	37	34_Fruit
35	35	36	37	38	35_Fruit
36	36	37	38	39	36_Fruit
37	37	38	39	40	37_Fruit
38	38	39	40	41	38_Fruit
39	39	40	41	42	39_Fruit
40	40	41	42	43	40_Fruit
41	41	42	43	44	41_Fruit
42	42	43	44	45	42_Fruit
43	43	44	45	46	43_Fruit
44	44	45	46	47	44_Fruit
45	45	46	47	48	45_Fruit
46	46	47	48	49	46_Fruit
47	47	48	49	50	47_Fruit
48	48	49	50	51	48_Fruit
49	49	50	51	52	49_Fruit
50	50	51	52	53	50_Fruit
51	51	52	53	54	51_Fruit
52	52	53	54	55	52_Fruit
53	53	54	55	56	53_Fruit
54	54	55	56	57	54_Fruit
55	55	56	57	58	55_Fruit
56	56	57	58	59	56_Fruit
57	57	58	59	60	57_Fruit
58	58	59	60	61	58_Fruit
59	59	60	61	62	59_Fruit
60	60	61	62	63	60_Fruit
61	61	62	63	64	61_Fruit
62	62	63	64	65	62_Fruit
63	63	64	65	66	63_Fruit
64	64	65	66	67	64_Fruit
65	65	66	67	68	65_Fruit
66	66	67	68	69	66_Fruit
67	67	68	69	70	67_Fruit
68	68	69	70	71	68_Fruit
69	69	70	71	72	69_Fruit
70	70	71	72	73	70_Fruit
71	71	72	73	74	71_Fruit
72	72	73	74	75	72_Fruit
73	73	74	75	76	73_Fruit
74	74	75	76	77	74_Fruit
75	75	76	77	78	75_Fruit
76	76	77	78	79	76_Fruit
77	77	78	79	80	77_Fruit
78	78	79	80	81	78_Fruit
79	79	80	81	82	79_Fruit
80	80	81	82	83	80_Fruit
81	81	82	83	84	81_Fruit
82	82	83	84	85	82_Fruit
83	83	84	85	86	83_Fruit
84	84	85	86	87	84_Fruit
85	85	86	87	88	85_Fruit
86	86	87	88	89	86_Fruit
87	87	88	89	90	87_Fruit
88	88	89	90	91	88_Fruit
89	89	90	91	92	89_Fruit
90	90	91	92	93	90_Fruit
91	91	92	93	94	91_Fruit
92	92	93	94	95	92_Fruit
93	93	94	95	96	93_Fruit
94	94	95	96	97	94_Fruit
95	95	96	97	98	95_Fruit
96	96	97	98	99	96_Fruit
97	97	98	99	100	97_Fruit
98	98	99	100	101	98_Fruit
99	99	100	101	102	99_Fruit
100	100	101	102	103	100_Fruit
101	101	102	103	104	101_Fruit
102	102	103	104	105	102_Fruit
103	103	104	105	106	103_Fruit
104	104	105	106	107	104_Fruit
105	105	106	107	108	105_Fruit
106	106	107	108	109	106_Fruit
107	107	108	109	110	107_Fruit
108	108	109	110	111	108_Fruit
109	109	110	111	112	109_Fruit
110	110	111	112	113	110_Fruit
111	111	112	113	114	111_Fruit
112	112	113	114	115	112_Fruit
113	113	114	115	116	113_Fruit
114	114	115	116	117	114_Fruit
115	115	116	117	118	115_Fruit
116	116	117	118	119	116_Fruit
117	117	118	119	120	117_Fruit
118	118	119	120	121	118_Fruit
119	119	120	121	122	119_Fruit
120	120	121	122	123	120_Fruit
121	121	122	123	124	121_Fruit
122	122	123	124	125	122_Fruit
123	123	124	125	126	123_Fruit
124	124	125	126	127	124_Fruit
125	125	126	127	128	125_Fruit
126	126	127	128	129	126_Fruit
127	127	128	129	130	127_Fruit
128	128	129	130	131	128_Fruit
129	129	130	131	132	129_Fruit
130	130	131	132	133	130_Fruit
131	131	132	133	134	131_Fruit
132	132	133	134	135	132_Fruit
133	133	134	135	136	133_Fruit
134	134	135	136	137	134_Fruit
135	135	136	137	138	135_Fruit
136	136	137	138	139	136_Fruit
137	137	138	139	140	137_Fruit
138	138	139	140	141	138_Fruit
139	139	140	141	142	139_Fruit
140	140	141	142	143	140_Fruit
141	141	142	143	144	141_Fruit
142	142	143	144	145	142_Fruit
143	143	144	145	146	143_Fruit
144	144	145	146	147	144_Fruit
145	145	146	147	148	145_Fruit
146	146	147	148	149	146_Fruit
147	147	148	149	150	147_Fruit
148	148	149	150	151	148_Fruit
149	149	150	151	152	149_Fruit
150	150	151	152	153	150_Fruit
151	151	152	153	154	151_Fruit
152	152	153	154	155	152_Fruit
153	153	154	155	156	153_Fruit
154	154	155	156	157	154_Fruit
155	155	156	157	158	155_Fruit
156	156	157	158	159	156_Fruit
157	157	158	159	160	157_Fruit
158	158	159	160	161	158_Fruit
159	159	160	161	162	159_Fruit
160	160	161	162	163	160_Fruit
161	161	162	163	164	161_Fruit
162	162	163	164	165	162_Fruit
163	163	164	165	166	163_Fruit
164	164	165	166	167	164_Fruit
165	165	166	167	168	165_Fruit
166	166	167	168	169	166_Fruit
167	167	168	169	170	167_Fruit
168	168	169	170	171	168_Fruit
169	169	170	171	172	169_Fruit
170	170	171	172	173	170_Fruit
171	171	172	173	174	171_Fruit
172	172	173	174	175	172_Fruit
173	173	174	175	176	173_Fruit
174	174	175	176	177	174_Fruit
175	175	176	177	178	175_Fruit
176	176	177	178	179	176_Fruit
177	177	178	179	180	177_Fruit
178	178	179	180	181	178_Fruit
179	179	180	181	182	179_Fruit
180	180	181	182	183	180_Fruit
181	181	182	183	184	181_Fruit
182	182	183	184	185	182_Fruit
183	183	184	185	186	183_Fruit
184	184	185	186	187	184_Fruit
185	185	186	187	188	185_Fruit
186	186	187	188	189	186_Fruit
187	187	188	189	190	187_Fruit
188	188	189	190	191	188_Fruit
189	189	190	191	192	189_Fruit
190	190	191	192	193	190_Fruit
191	191	192	193	194	191_Fruit
192	192	193	194	195	192_Fruit
193	193	194	195	196	193_Fruit
194	194	195	196	197	194_Fruit
195	195	196	197	198	195_Fruit
196	196	197	198	199	196_Fruit
197	197	198	199	200	197_Fruit
198	198	199	200	201	198_Fruit
199	199	200	201	202	199_Fruit
200	200	201	202	203	200_Fruit
201	201	202	203	204	201_Fruit
202	202	203	204	205	202_Fruit
203	203	204	205	206	203_Fruit
204	204	205	206	207	204_Fruit
205	205	206	207	208	205_Fruit
206	206	207	208	209	206_Fruit
207	207	208	209	210	207_Fruit
208	208	209	210	211	208_Fruit
209	209	210	211	212	209_Fruit
210	210	211	212	213	210_Fruit
211	211	212	213	214	

Step 6: Evaluate the model's performance. Once the process finishes, Weka will display various performance measures like accuracy, precision, recall, and ROC curve (for classification tasks) or RMSE and MAE (for regression tasks). These measures can be found in the "Result list" on the right side of the window.

Step 7: Analyse the results and interpret them. Examine the performance measures to assess the model's quality and suitability for your dataset. Compare different models or validation methods if you have tried more than one.

Step 8: Repeat steps 4-7 with different algorithms or validation methods if desired. This will help you compare the performance of different models and choose the best one.

Result:

Thus, the simple data validation and testing dataset using Weka was implemented.

EX NO: 03

PLAN THE ARCHITECTURE FOR REAL TIME APPLICATION

DATE:

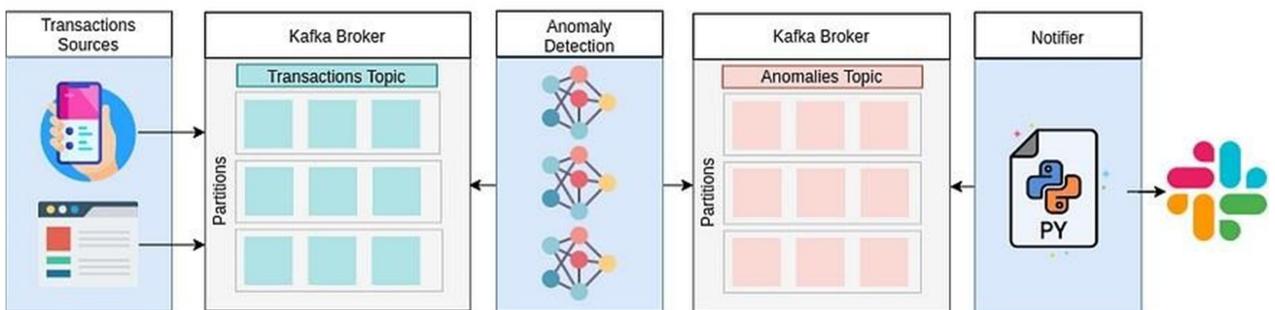
Aim:

To make real-time predictions with incoming stream data from Apache Kafka, and to implement notification messages for credit card transactions, GPS logs, system consumption metrics.

Project ideas:

- Train an anomaly detection algorithm using unsupervised machine learning.
- Create a new data producer that sends the transactions to a Kafka topic.
- Read the data from the Kafka topic to make the prediction using the trained ml model.
- If the model detects that the transaction is not an inlier, send it to another Kafka topic.
- Create the last consumer that reads the anomalies and sends an alert to a Slack channel.

Architecture:



Procedure:

Step 1: Project Structure:

- i) First, Check the Settings.Py; It Has Some Variables to Set, Like the Kafka Broker Host and Port; Leave the Ones by Default (Listening On Localhost And Default Ports Of Kafka And Zookeeper).
- ii) The Streaming/Utils.Py File Contains the Configurations to Create Kafka Consumers and Producers.
- iii) Install The Requirements

Step 2: Train the Model

- i) To generate random data; it will have two variables
- ii) Isolation Forest model to detect the outliers; (To isolate the data points by

tracing random lines over one of the (sampled) variables' axes and, after several iterations, measure how "hard" was to isolate each observation).

Step 3: Create the Topics

- i)** "transactions," where the producer will send new transaction records.
- ii)** "anomalies," the module that detects anomalies will send the data, and the last consumer will read it to send a slack notification:

Step 4: Transaction Producer

- i)** To generate the first producer that will send new data to the Kafka topic "transactions"; use the confluent-Kafka package; in the file streaming/producer.py.
- ii)** Producer will send data to a Kafka topic, with a probability of OUTLIERS_GENERATION_PROBABILITY; the data will come from an "outlier generator," will send an auto-incremental id, the data needed for the machine learning model and the current time in UTC.

Step 5: Outlier Detector Consumer

- i)** To make the predictions, and filter the outliers. Done in the streaming/anomalies_detector.py file
- ii)** The consumer read messages from the "transactions" topic and a consumer sent outliers to the "anomalies".

Step 6: Slack notification

- i)** To take some actions with these detected outliers; in a real-life scenario, it could block a transaction, scale a server, generate a recommendation, send an alert to an administrative user.

Program:

Step1: pip install -r requirements.txt.

Step2: Train and build the model

```
from joblib import dump
import numpy as np
from sklearn.ensemble import IsolationForest

rng = np.random.RandomState(42)
# Generate train data
X = 0.3 * rng.randn(500, 2)
X_train = np.r_[X + 2, X - 2]
X_train = np.round(X_train, 3)

# fit the model
clf = IsolationForest(n_estimators=50, max_samples=500, random_state=rng, contamination=0.01)
clf.fit(X_train)

dump(clf, './isolation_forest.joblib')
```

Step 3: Create the topics

```
kafka-topics.sh --zookeeper localhost:2181 --topic transactions --create --partitions 3 --replication-factor 1
kafka-topics.sh --zookeeper localhost:2181 --topic anomalies --create --partitions 3 --replication-factor 1
```

Step 4: Transaction producer

```
import json
import random
import time
from datetime import datetime
import numpy as np
from settings import TRANSACTIONS_TOPIC, DELAY, OUTLIERS_GENERATION_PROBABILITY
from streaming.utils import create_producer
```

```

_id = 0
producer = create_producer()

if producer is not None:
    while True:
        # Generate some abnormal observations
        if random.random() <= OUTLIERS_GENERATION_PROBABILITY:
            X_test = np.random.uniform(low=-4, high=4, size=(1, 2)).tolist()
        else:
            X = 0.3 * np.random.randn(1, 2)
            X_test = (X + np.random.choice(a=[2, -2], size=1, p=[0.5, 0.5])).tolist()
            X_test = np.round(X_test, 3).tolist()

        current_time = datetime.utcnow().isoformat()

        record = {"id": _id, "data": X_test, "current_time": current_time}
        record = json.dumps(record).encode("utf-8")

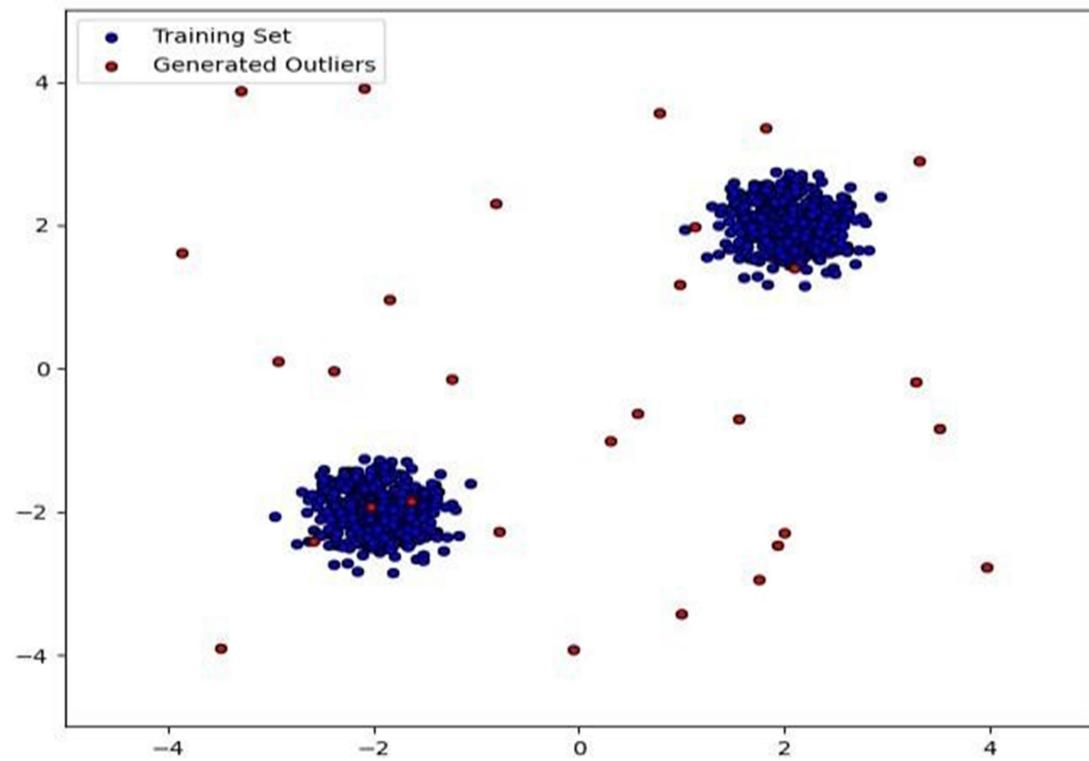
        producer.produce(topic=TRANSACTIONS_TOPIC,
                          value=record,
                          key=str(_id))
        producer.flush()
        _id += 1
        time.sleep(DELAY)

```

kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic transactions

Output

Anomaly detection



Step 5: Outlier Detector Consumer.

```
import json
import os
from joblib import load
import logging
from multiprocessing import Process
import numpy as np
from streaming.utils import create_producer, create_consumer
from settings import (TRANSACTIONS_TOPIC, TRANSACTIONS_CONSUMER_GROUP,
                      ANOMALIES_TOPIC, NUM_PARTITIONS)

model_path = os.path.abspath('../model/isolation_forest.joblib')
clf = load(model_path)

def detect():
    consumer = create_consumer(topic=TRANSACTIONS_TOPIC,
                                group_id=TRANSACTIONS_CONSUMER_GROUP)

    producer = create_producer()

    while True:
        message = consumer.poll(timeout=50)
        if message is None:
            continue
        if message.error():
            logging.error("Consumer error: {}".format(message.error()))
            continue

        # Message that came from producer
        record = json.loads(message.value().decode('utf-8'))
        data = record["data"]

        prediction = clf.predict(data)

        # If an anomaly comes in, send it to anomalies topic
        if prediction[0] == -1:
            score = clf.score_samples(data)
            record["score"] = np.round(score, 3).tolist()

            _id = str(record["id"])
            record = json.dumps(record).encode("utf-8")

            producer.produce(topic=ANOMALIES_TOPIC,
                              value=record,
                              key=_id)
            producer.flush()

    consumer.close()

    # One consumer per partition
    for _ in range(NUM_PARTITIONS):
        p = Process(target=detect)
        p.start()
```

CHATBOT ALERT NOTIFICATION



rodrigo-arenas 1:19 PM
new



anomalies-alerts APP 1:19 PM

```
{"id": 136, "data": [[3.052, 2.606]], "current_time": "2021-06-17T18:19:55.734364", "score": [-0.691]}

{"id": 183, "data": [[2.973, 2.7]], "current_time": "2021-06-17T18:20:00.743853", "score": [-0.698]}

{"id": 190, "data": [[2.569, 1.276]], "current_time": "2021-06-17T18:20:01.488960", "score": [-0.647]}

{"id": 195, "data": [[2.515, 1.145]], "current_time": "2021-06-17T18:20:02.021038", "score": [-0.634]}

{"id": 197, "data": [[3.096, 3.062]], "current_time": "2021-06-17T18:20:02.232794", "score": [-0.745]}

 {"id": 216, "data": [[3.74, 2.731]], "current_time": "2021-06-17T18:20:04.261914", "score": [-0.704]}

 {"id": 231, "data": [[2.658, 3.264]], "current_time": "2021-06-17T18:20:05.860120", "score": [-0.711]}
```

Send a message to **anomalies-alerts**

```
{"id": 183, "data": [[2.973, 2.7]], "current_time": "2021-06-17T18:20:00.743853", "score": [-0.698]}

 {"id": 190, "data": [[2.569, 1.276]], "current_time": "2021-06-17T18:20:01.488960", "score": [-0.647]}

 {"id": 195, "data": [[2.515, 1.145]], "current_time": "2021-06-17T18:20:02.021038", "score": [-0.634]}

 {"id": 197, "data": [[3.096, 3.062]], "current_time": "2021-06-17T18:20:02.232794", "score": [-0.745]}

 {"id": 216, "data": [[3.74, 2.731]], "current_time": "2021-06-17T18:20:04.261914", "score": [-0.704]}

 1:20 {"id": 231, "data": [[2.658, 3.264]], "current_time": "2021-06-17T18:20:05.860120", "score": [-0.711]}

 {"id": 276, "data": [[3.23, 2.997]], "current_time": "2021-06-17T18:20:10.650696", "score": [-0.732]}

 {"id": 344, "data": [[2.92, 3.255]], "current_time": "2021-06-17T18:20:17.888443", "score": [-0.751]}

 {"id": 352, "data": [[3.483, 2.8]], "current_time": "2021-06-17T18:20:18.739169", "score": [-0.717]}
```

Send a message to **anomalies-alerts**

Step6: Slack notification

```
import logging
from slack import WebClient
from slack.errors import SlackApiError

from settings import (SLACK_API_TOKEN, SLACK_CHANNEL, ANOMALIES_TOPIC,
                      ANOMALIES_CONSUMER_GROUP)
from streaming.utils import create_consumer

client = WebClient(token=SLACK_API_TOKEN)
consumer = create_consumer(topic=ANOMALIES_TOPIC,
                           group_id=ANOMALIES_CONSUMER_GROUP)

while True:
    message = consumer.poll()
    if message is None:
        continue
    if message.error():
        logging.error("Consumer error: {}".format(message.error()))
        continue

    # Message that came from producer
    record = message.value().decode('utf-8')

    try:
        # Send message to slack channel
        response = client.chat_postMessage(channel=SLACK_CHANNEL,
                                             text=record)
    except SlackApiError as e:

        print(e.response["error"])

    consumer.commit()
```

Result:

Thus, the prediction of Real-time anomaly detection with Apache Kafka and Python has executes the streaming/Bot alerts notification.

EX NO: 04 (A)

DATE:

QUERY FOR STAR SCHEMA USING SQL SERVER MANAGEMENT STUDIO

Aim:

To execute and verify query for star schema using SQL Server Management Studio.

Procedure:

Step 1: Install SQLEXPR and SQLManagementStudio

Step 2: Launch SQL Server Management Studio

Step 3: Create new database and write query for creating Star schema table

Step 4: Execute the query for schema

Step 5: Explore the database diagram for Star schema

Query for Star Schema

USE DemoGO

CREATE TABLE DimProduct

(ProductKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,

ProductAltKey nvarchar(10) NOT NULL,

ProductName nvarchar(50) NULL, ProductDescription
nvarchar(100) NULL, ProductCategoryName nvarchar(50))

GO

CREATE TABLE DimCustomer

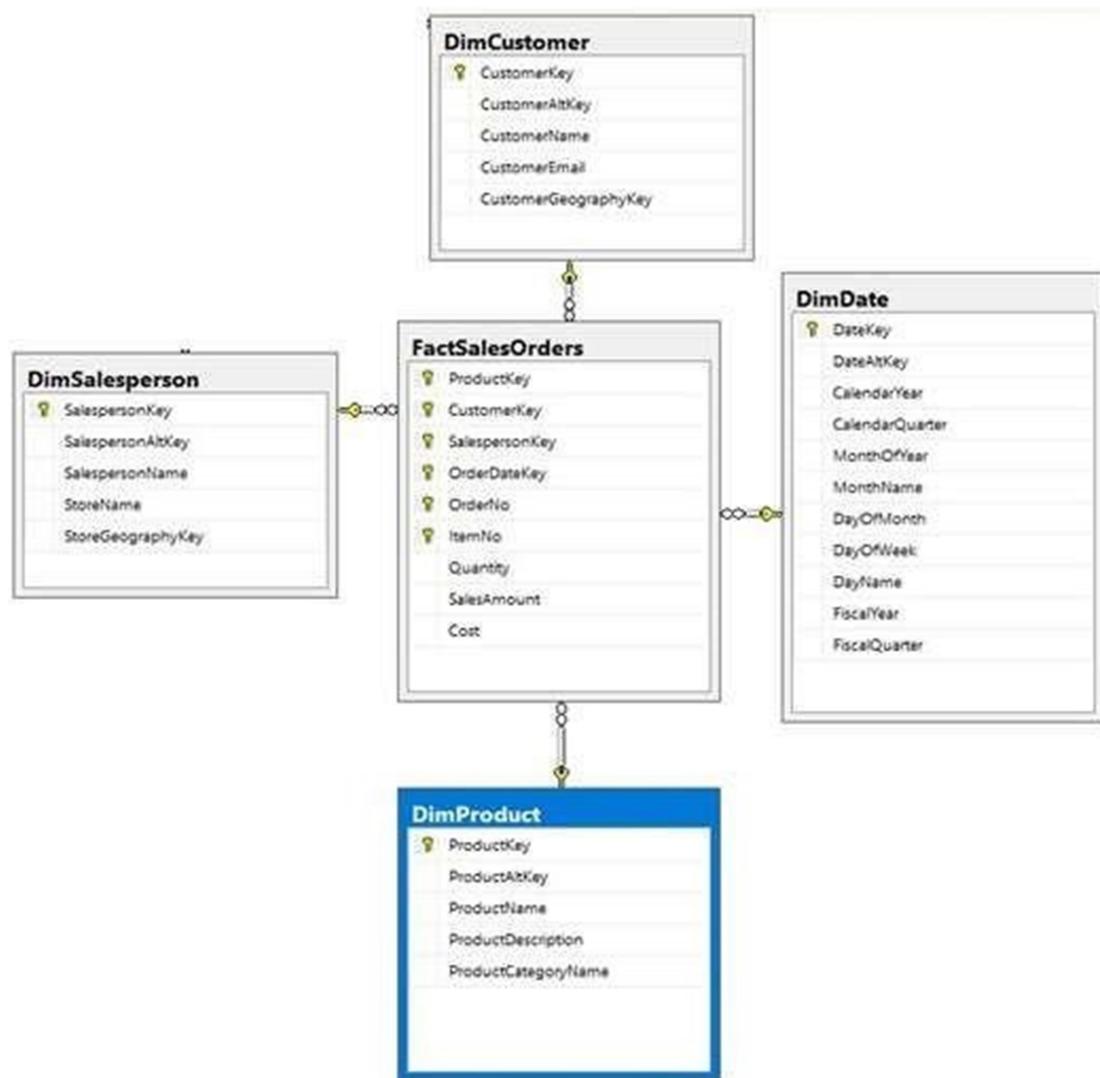
(CustomerKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,

CustomerAltKey nvarchar(10) NOT NULL, CustomerName nvarchar(50) NULL,
CustomerEmail nvarchar(100) NULL,

CustomerGeographyKey int NULL)

GO

Output



```
CREATE TABLE DimSalesperson
(SalespersonKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,
SalespersonAltKey nvarchar(10) NOT NULL,
SalespersonName nvarchar(50) NULL, StoreName nvarchar(50) NULL,
SalespersonGeographyKey int NULL) GO
```

```
CREATE TABLE DimDate
```

```
(DateKey int NOT NULL PRIMARY KEY NONCLUSTERED, DateAltKey
datetime NOT
NULL, CalendarYear int NOT NULL, CalendarQuarter int NOT NULL,
MonthOfYear int NOT NULL, [MonthName]nvarchar(15) NOT NULL,
[DayOfMonth]int NOT NULL,
[DayOfWeek]int NOT NULL, [DayName]nvarchar(15) NOT NULL,
FiscalYear int NOT NULL, FiscalQuarter int NOT NULL)
```

```
GO
```

```
CREATE TABLE FactSalesOrders
```

```
(ProductKey int NOT NULL REFERENCES DimProduct(ProductKey),
CustomerKey int NOT NULL REFERENCES DimCustomer(CustomerKey),
SalespersonKey int NOT NULL REFERENCES
DimSalesperson(SalespersonKey), OrderDateKey int NOT NULL
REFERENCES DimDate(DateKey),
OrderNo int NOT NULL, ItemNo int NOT NULL, Quantity int NOT NULL,
SalesAmount money NOT NULL,
Cost money NOT NULL
```

Result:

Thus, the Query for Star Schema was created and executed successfully.

EX NO: 04 (B)

DATE:

QUERY FOR SNOWFLAKE SCHEMA USING SQL SERVER MANAGEMENT STUDIO

Aim:

To execute and verify query for Snowflake schema using SQL Server Management Studio.

Procedure:

Step 1: Install SQLEXPR and SQLManagementStudio

Step 2: LaunchSQL Server Management Studio

Step 3: Create new database and write query for creating Star schema table

Step 4: Execute the query

Step 5: Explore the database diagram for SnowFlake schema

Step 6: Connect the Geography table with Salesperson & Product Geography key

Query for SnowFlake Schema

USE Demo

GO

CREATE TABLE DimProduct

(ProductKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,
ProductAltKey nvarchar(10) NOT NULL,

ProductName nvarchar(50) NULL, ProductDescription
nvarchar(100) NULL, ProductCategoryName nvarchar(50))

GO

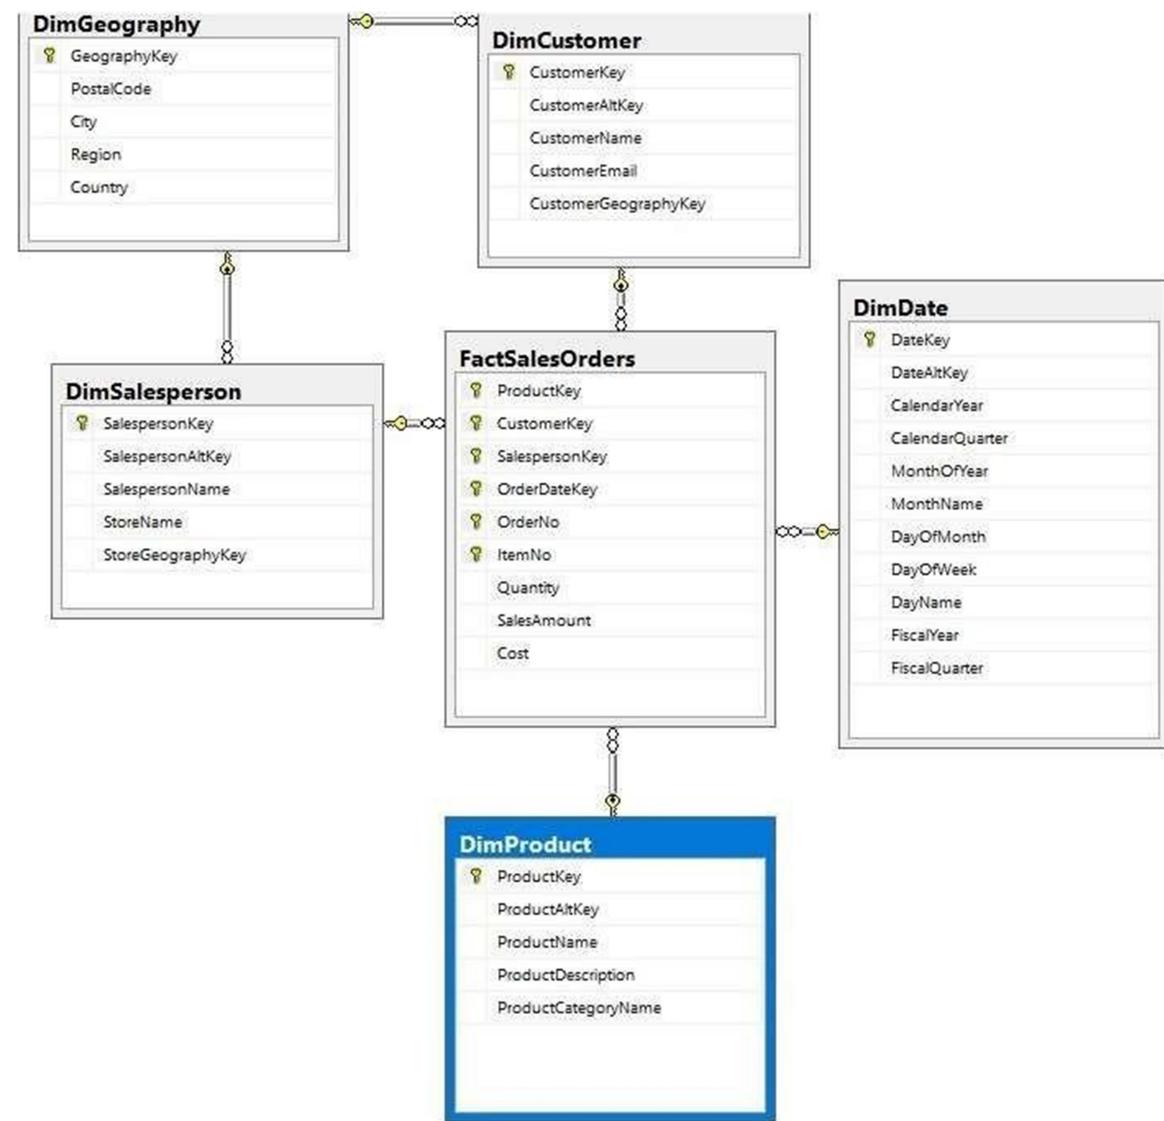
CREATE TABLE DimCustomer

(CustomerKey int identity NOT NULL PRIMARY KEY
NONCLUSTERED,

CustomerAltKey nvarchar(10) NOT NULL,

```
CustomerName nvarchar(50) NULL, CustomerEmail nvarchar(100) NULL,  
CustomerGeographyKey int NULL) GO
```

Output



```

CREATE TABLE DimSalesperson
(SalespersonKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,
SalespersonAltKey nvarchar(10) NOT NULL,
SalespersonName nvarchar(50) NULL, StoreName nvarchar(50) NULL,
SalespersonGeographyKey int NULL) GO

CREATE TABLE DimGeography
(GeographyKey int identity NOT NULL PRIMARY KEY NONCLUSTERED,
PostalCode nvarchar(15) NULL,
City nvarchar(50) NULL, Region nvarchar(50) NULL, Country nvarchar(50)
NULL) GO

CREATE TABLE FactSalesOrders
(ProductKey int NOT NULL REFERENCES DimProduct(ProductKey),
CustomerKey int NOT NULL REFERENCES DimCustomer(CustomerKey),
SalespersonKey int NOT NULL
REFERENCES DimSalesperson(SalespersonKey), OrderNo int NOT NULL,
ItemNo int NOT NULL, Quantity int NOT NULL,
SalesAmount money NOT NULL, Cost money NOT NULL
CONSTRAINT[PK_FactSalesOrders] PRIMARY KEY NONCLUSTERED (
[ProductKey],[CustomerKey],[SalesPersonKey],[OrderNo], [ItemNo]
))

```

Result:

Thus, the Query for Snowflake Schema was created and executed successfully.

EX NO: 05

DATE:

DESIGN DATA WAREHOUSE FOR REAL TIME APPLICATIONS

Aim:

To design and execute data warehouse for real time application using SQL Server Management Studio.

Procedure:

Step 1: Launch SQL Server Management Studio

Step 2: Explore the created database

Step 3: **3.1** Right-click on the table name and click on the Edit top 200 rows option.

3.2. Enter the data inside the table or use the top 1000 rows option and enter the query.

Step 4: Execute the query, and the data will be updated in the table.

Step 5: Right-click on the database and click on the tasks option. Use the import data option to import files to the database.

Output

The screenshot shows the Object Explorer on the left and a query results window on the right. The query window displays the following SQL statement:

```
select * from dbo.person
```

The results show the following data:

person_id	first_name	last_name	gender
1	John	Doe	M
2	Jane	Doe	F
3	Kevi	B	M
4	Nia	V	F
5	Normal	B	M
6	Kevaya	M	F

A message at the bottom of the results window says "Query executed successfully."

Import CSV file

The screenshot shows the Object Explorer on the left and a query results window on the right. The query window displays the following SQL statement:

```
select * from dbo.[csv_result-airline]
```

The results show the following data:

id	passenger_number	Date
19	170	1950-07-01
20	170	1950-08-01
21	158	1950-09-01
22	133	1950-10-01
23	114	1950-11-01
24	140	1950-12-01
25	145	1951-01-01
26	150	1951-02-01
27	178	1951-03-01
28	163	1951-04-01
29	172	1951-05-01
30	178	1951-06-01
31	199	1951-07-01
32	199	1951-08-01
33	184	1951-09-01

A message at the bottom of the results window says "Query executed successfully."

Sample Query

```
INSERT INTO dbo.person(first_name,last_name,gender) VALUES  
('Kavi','S','M'), ('Nila','V','A'), ('Nirmal','B','M'), ('Kaviya','M','F');
```

```
SELECT * FROM dbo.person
```

Result:

Thus, the data warehouse for real-time applications was designed successfully.

EX NO: 06

DATE:

ANALYSE THE DIMENSIONAL MODELING

Aim:

To implement the creation of table dimensions and analysis of data model.

Procedure:

Step1:Identify the Business Process

Step2:Identify the Grain

Step3:Identify the Dimensions

Step4:Identify the Facts

Step5:Build the Schema

Implementation:

Create the data

warehouse

create database

TopHireDW go

use TopHireDW

go

-- Create Date Dimension

if exists (select * from sys.tables

where name = 'DimDate') drop table

DimDate go

create table DimDate

(DateKey int not null primary key,

[Year] varchar(7), [Month] varchar(7), [Date] date,

DateString varchar(10)) go

-- Populate

Date

Dimension

truncate

table

DimDate go

declare @i int, @Date date, @StartDate date, @EndDate date, @DateKey int,

@DateString varchar(10), @Year

varchar(4), @Month varchar(7),

@Date1 varchar(20)

set @StartDate

= '2006-01-01'

set @EndDate =

```
'2016-12-31' set
```

```
@Date =
```

```
insert into DimDate (DateKey, [Year], [Month], [Date], DateString) values (0,  
'Unknown', 'Unknown', '0001-01-01', 'Unknown') --The unknown row while @Date  
<= @EndDate
```

```
begin set @DateString = convert(varchar(10), @Date, 20) set @DateKey =  
convert(int, replace(@DateString,'-','')) set @Year = left(@DateString,4) set  
@Month = left(@DateString, 7)
```

```
insert into DimDate (DateKey, [Year], [Month], [Date], DateString) values  
(@DateKey, @Year, @Month, @Date, @DateString) set @Date = dateadd(d, 1,  
@Date) end go
```

```
select * from DimDate
```

```
-- Create Customer dimension
```

```
if exists (select * from sys.tables where name =  
'DimCustomer') drop table DimCustomer go
```

```
create table DimCustomer
```

```
( CustomerKey int not null identity(1,1) primary key,  
CustomerId varchar(20) not null,  
CustomerName varchar(30), DateOfBirth date, Town varchar(50),  
TelephoneNo varchar(30), DrivingLicenceNo varchar(30), Occupation  
varchar(30)
```

```
)
```

```
go
```

```
insert into DimCustomer (CustomerId, CustomerName, DateOfBirth, Town,  
TelephoneNo, DrivingLicenceNo, Occupation) select * from  
HireBase.dbo.Customer select * from DimCustomer -- Create Van  
dimension
```

```
if exists (select * from sys.tables where name =  
'DimVan') drop table DimVan go
```

```
create table DimVan
```

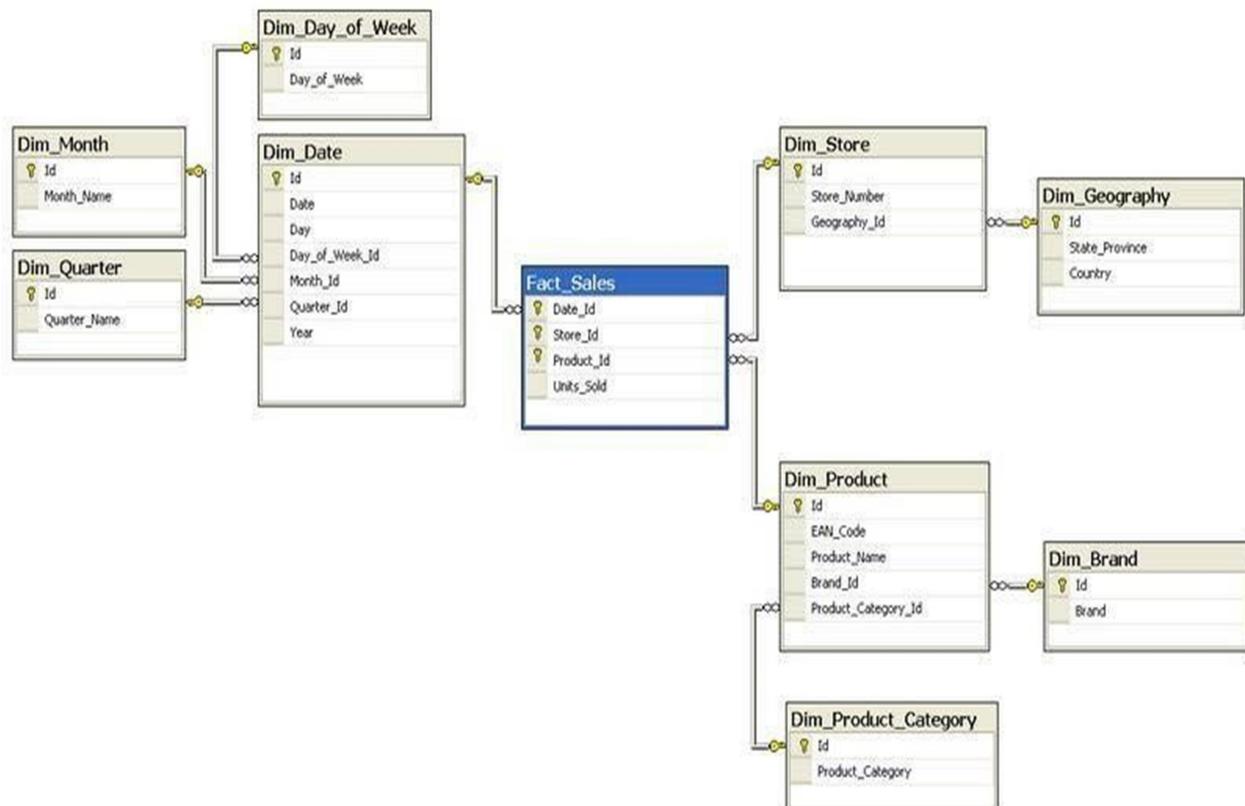
```
( VanKey int not null identity(1,1) primary key,  
RegNo varchar(10) not null,
```

```
Make varchar(30), Model varchar(30), [Year] varchar(4),
```

Colour varchar(20), CC int, Class varchar(10)

Output

Snowflake Schema image source of dimensional data model



```

)
go
insert into DimVan (RegNo, Make, Model, [Year], Colour,
CC, Class) select * from HireBase.dbo.Van go select *
from DimVan -- Create Hire fact table

if exists (select * from sys.tables where name =
'FactHire') drop table FactHire go
create table FactHire

( SnapshotDateKey int not null, --Daily periodic snapshot fact table
HireDateKey int not null, CustomerKey int not null, VanKey int not null, --
Dimension Keys
HireId varchar(10) not null, --Degenerate Dimension
NoOfDay int, VanHire money, SatNavHire money,
Insurance money, DamageWaiver money, TotalBill money
)
go
select * from FactHire
tool with MySQL query

```

Result:

Thus, the implementation and analysed of data dimension model using weka implemented successfully.

EX NO: 07	CASE STUDY USING OLAP
DATE:	

Aim:

To evaluate the implementation and impact of OLAP technology in a real-world business context, analysing its effectiveness in enhancing data analysis, decision-making, and overall operational efficiency.

Introduction:

OLAP stands for **On-Line Analytical Processing**. OLAP is a classification of software technology which authorizes analysts, managers, and executives to gain insight into information through fast, consistent, interactive access in a wide variety of possible views of data that has been transformed from raw information to reflect the real dimensionality of the enterprise as understood by the clients .It is used to analyse business data from different points of view. Organizations collect and store data from multiple data sources, such as websites, applications, smart meters, and internal systems.

Methodology

OLAP (Online Analytical Processing) methodology refers to the approach and techniques used to design, create, and use OLAP systems for efficient multidimensional data analysis. Here are the key components and steps involved in the OLAP methodology:

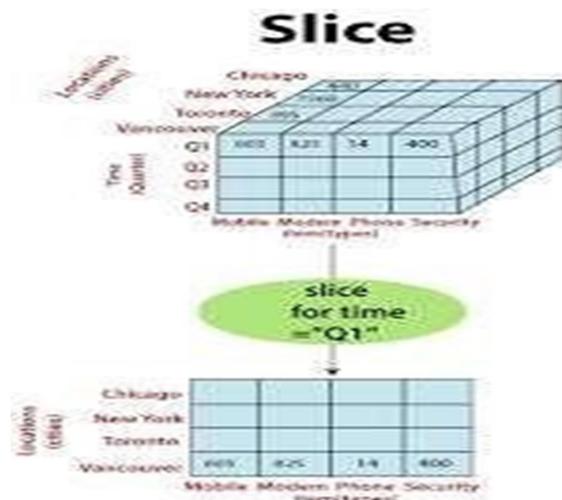
1. Requirement Analysis:

The process begins with understanding the specific analytical requirements of the users. Analysts and stakeholders define the dimensions, measures, hierarchies, and data sources that will be part of the OLAP system. This step is crucial to ensure that the OLAP system meets the business needs.

2. Dimensional Modeling:

Dimension tables are designed to represent attributes like time, geography, and product categories. Fact tables contain the numerical data (measures) and the keys to dimension tables.

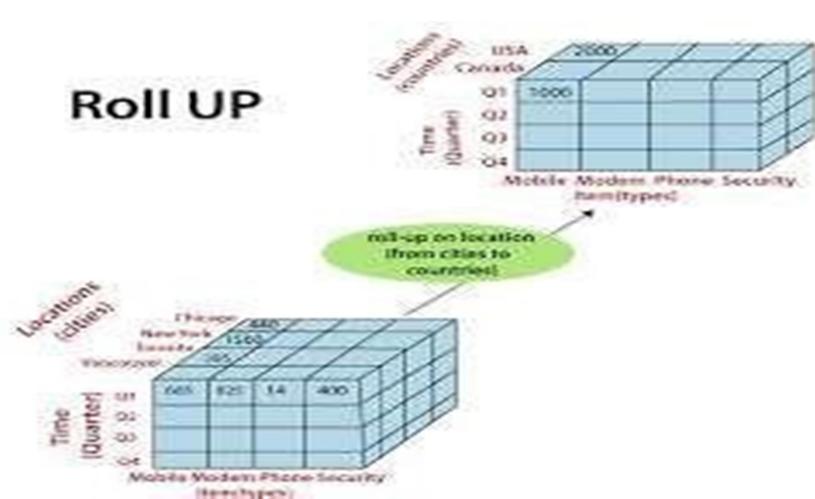
Slice



Dice



Roll Up



3. Star Schema:

This is a common design in OLAP systems where the fact table is at the center, connected to dimension tables.

4. Data Extraction and Transformation:

Data is extracted from various sources, cleaned, and transformed into a format suitable for OLAP analysis. This may involve data aggregation, cleansing, and integration.

5. Data Loading:

The prepared data is loaded into the OLAP database or cube. This step includes populating the dimension and fact tables and creating the data cube structure.

Operations in OLAP

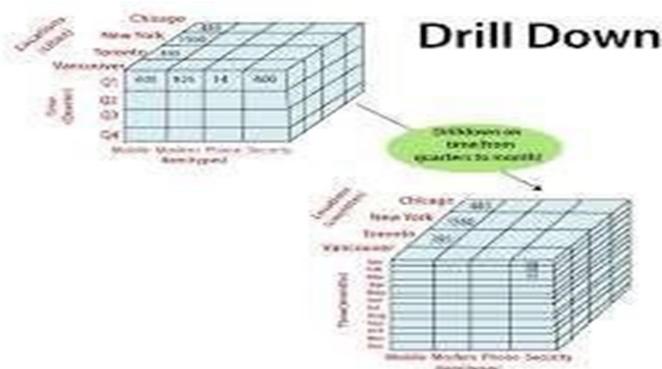
In OLAP (Online Analytical Processing), operations are the fundamental actions performed on multidimensional data cubes to retrieve, analyze, and present data in a way that facilitates decision-making and data exploration. The main operations in OLAP are:

1. **Slice:** Slicing involves selecting a single dimension from a multidimensional cube to view a specific "slice" of the data. For example, you can slice the cube to view sales data for a particular month, product category, or region.
2. **Dice:** Dicing is the process of selecting specific values from two or more dimensions to create a subcube. It allows you to focus on a particular combination of attributes. For example, you can dice the cube to view sales data for a specific product category and region within a certain time frame.
3. **Roll-up (Drill-up):** Roll-up allows you to move from a more detailed level of data to a higher-level summary. For instance, you can roll up from daily sales data to monthly or yearly sales data, aggregating the information.
4. **Drill-down (Drill-through):** Drill-down is the opposite of roll-up, where you move from a higher-level summary to a more detailed view of the data. For example, you can drill down from yearly sales data to quarterly, monthly, and daily data, getting more granularity.

Pivot



Drill Down



5. **Pivot (Rotate):** Pivoting involves changing the orientation of the cube, which means swapping dimensions to view the data from a different perspective. This operation is useful for exploring data in various ways.
6. **Slice and Dice:** Combining slicing and dicing allows you to select specific values from different dimensions to create sub cubes. This operation helps you focus on a highly specific subset of the data.
7. **Drill-across:** Drill-across involves navigating between cubes that are related but have different dimensions or hierarchies. It allows users to explore data across different OLAP cubes.
8. **Data Filtering:** In OLAP, you can filter data to view only specific data points or subsets that meet certain criteria. This operation is useful for narrowing down data to what is most relevant for analysis.

Real time example

One of the real time examples of olap is Market Basket Analysis. Let us discuss in detail about the example

Market Basket Analysis:

- A data mining technique, is typically performed using algorithms like Apriori, FP- growth, or Eclat. These algorithms are designed to discover associations or patterns in transaction data, such as retail sales.
- While traditional OLAP (Online Analytical Processing) is not the primary tool for market basket analysis, it can play a supporting role. Here's how OLAP can complement market basket analysis in more detail:

1. Data Integration:

Gather and integrate transaction data from various sources, such as point-of-sale systems, e-commerce platforms, or other transactional databases. Clean and pre-process the data, ensuring that it is in a format suitable for analysis.

2. Data Modeling:

Design a data model that will be used in the OLAP cube. In the context of market basket analysis, consider the following dimensions and measures:

Dimensions:

Time (e.g., day, week, month)

Products (individual items or product categories) Customers (if you want to analyse customer behaviour).

Measures:

- The count of transactions containing specific items or itemsets.
- The count of products in each transaction.
- Any other relevant metrics, such as revenue, quantity, or profit.

3. Data Loading:

Load the integrated and preprocessed transaction data into the OLAP cube. Ensure that the cube is regularly updated to reflect the most recent data.

4. OLAP Cube Design:

Define hierarchies and relationships within the cube to enable effective analysis. For instance, you might have hierarchies that allow drilling down from product categories to individual products.

5. Market Basket Analysis:

Although OLAP cubes are not designed for direct market basket analysis, they can facilitate it in several ways.

Conclusion

OLAP is a powerful technology for businesses and organizations seeking data insights, informed decisions, and performance improvement. It enables multidimensional data analysis, especially in complex, data-intensive environments. It is a crucial technology for organizations seeking to gain insights from their data and make informed decisions. It empowers businesses to analyse data efficiently and effectively, offering a competitive advantage in today's data-driven world.

EX.NO: 08	CASE STUDY USING OLTP
DATE:	

Aim:

Develop an OLTP system that enables the e-commerce company to process a high volume of online orders, track inventory, manage customer information, and handle financial transactions in real-time, ensuring data integrity and providing a seamless shopping experience for customers.

Introduction:

In today's digital age, businesses across various industries are relying heavily on technology to streamline their operations and provide seamless services to their customers. One crucial aspect of this technological transformation is the development and implementation of efficient Online Transaction Processing (OLTP) systems. This case study delves into the design and implementation of an OLTP system for a fictional e-commerce company, "TechTrend Electronics," and examines the key considerations, challenges, and aims associated with such a project. This case study aims to showcase the process of developing an OLTP system tailored to TechTrend Electronics' unique requirements. The objective is to ensure that the company can efficiently handle a multitude of real-time transactions while maintaining data accuracy and providing a seamless shopping experience for its customers.

Methodology:

The methodology for developing an OLTP (Online Transaction Processing) system for a case study involves a systematic approach to designing, implementing, and testing the system. Below is a step-by-step methodology for creating an OLTP system for a case study, using the fictional e-commerce company "Tech Trend Electronics" as an example:

1. Database Design:

Develop a well-structured relational database schema that aligns with the business requirements. Normalize the data to eliminate redundancy and ensure data consistency. Create entity-relationship diagrams and define data models for key entities like customers, products, orders, payments, and inventory.

2. Technology Selection:

Choose appropriate technologies for the database management system (e.g., MySQL, PostgreSQL, Oracle) and programming languages (e.g., Java, Python, C#) for the OLTP system. Evaluate and select suitable frameworks, libraries, and tools that align with the chosen technologies.

3. System Architecture:

Design the system's architecture, which may include multiple application layers, a web interface, and a database layer. Implement a layered architecture, separating concerns for scalability, maintainability, and security.

4. User Authentication and Authorization:

Implement user authentication mechanisms to secure access to the system for both customers and staff. Define access control policies and user roles (e.g., customers, administrators, and employees) based on the principle of least privilege.

5. Transaction Processing Logic:

Develop the transaction processing logic, including handling order placement, inventory management, and payment processing in real-time. Ensure that transactions adhere to the ACID properties for data integrity.

6. Security Measures:

Implement security measures to protect customer data, financial information, and the system itself. Use encryption for sensitive data and ensure that the system is protected against common security threats (e.g., SQL injection, cross-site scripting).

7. Deployment and Monitoring:

Deploy the OLTP system in a production environment.

Implement monitoring tools to track system performance, identify bottlenecks, and generate reports for system administrators.

8. Maintenance and Updates:

Establish a plan for system maintenance and regular updates to address issues, enhance functionality, and adapt to changing business needs.

Real World Example

In a real-world scenario, let's consider an e-commerce platform as an example of an OLTP system. The platform processes millions of transactions every day.

Here's a breakdown of how the system functions: Users can browse through the website, add products to their carts, and complete the checkout process. As a user completes the checkout process, a new transaction is created. This transaction contains information about the products purchased, the buyer's details, the shipping address, and other relevant data. The system generates an invoice for the buyer and sends it via email. The system generates transaction reports, such as daily sales summaries or sales by product category, for internal use and management. In this scenario, the e-commerce platform acts as an OLTP system, with its transaction processing capabilities and the real-time updates to inventory and order details being key components. Here's an alternative approach using OLAP: Aggregate sales data across all time and geographical locations, making it available for reporting and analysis.

Allow business managers to run complex analytical queries on this data, such as calculating average sales by product category, comparing sales trends between different regions, or identifying top-performing sales channels. Use OLAP tools like data warehouses and data cubes to enable fast, real-time access to aggregated data and to simplify the process of running complex analytical queries.

By leveraging OLAP capabilities, businesses can gain insights into their sales performance, identify trends and patterns, and make data-driven decisions. This can ultimately lead to increased revenue, better customer service, and more efficient use of resources.

Conclusion:

In conclusion, OLTP systems play a pivotal role in modern business operations, facilitating real-time transaction processing, data integrity, and customer interactions. These systems are designed for high concurrency, low-latency, and consistent data access, making them essential for day-to-day operations in various industries, such as finance, e-commerce, healthcare, and more.

Overall, OLTP systems are the backbone of modern business operations, ensuring the seamless execution of day-to-day transactions and delivering a positive customer experience.

EX NO: 09

DATE:

IMPLEMENTATION OF WAREHOUSE TESTING.

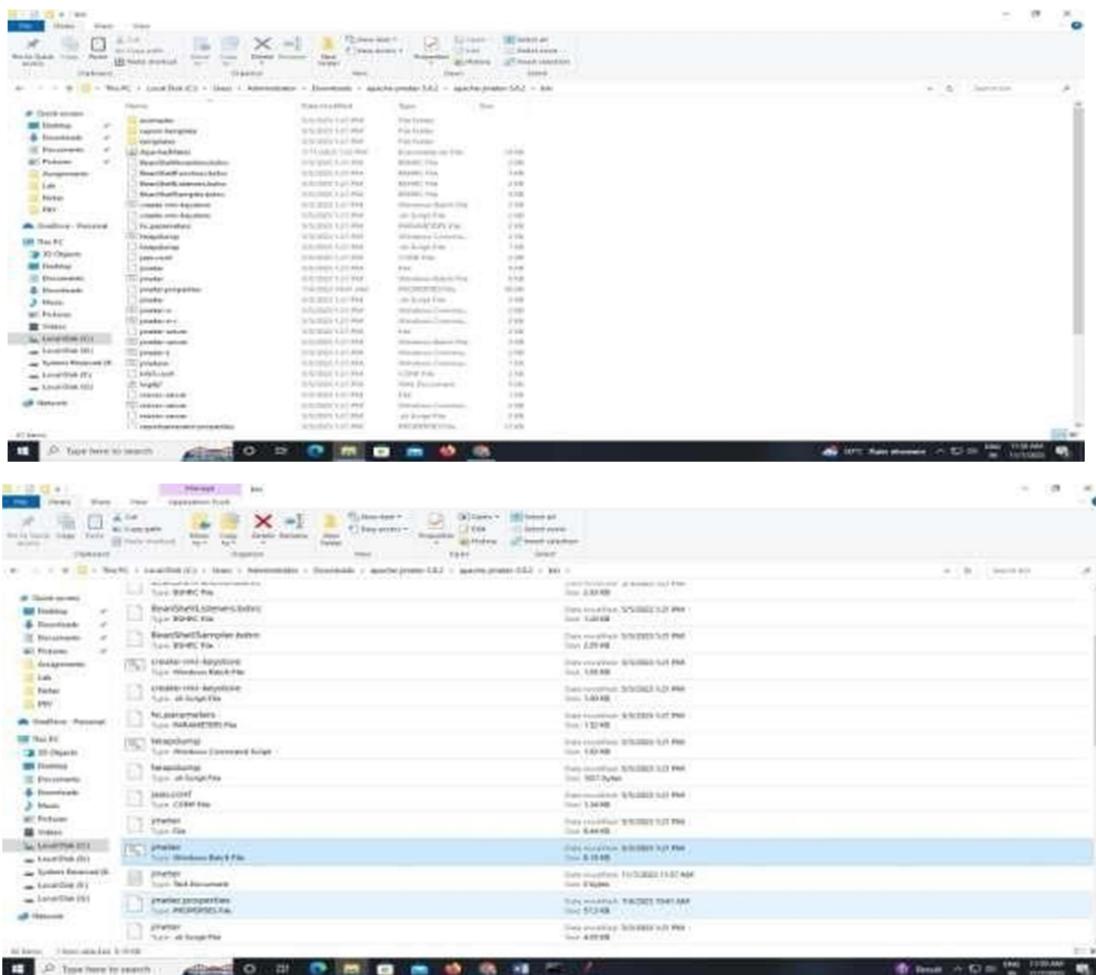
Aim:

To perform load testing using JMeter and interact with a SQL Server database using SQL Management Studio, you'll need to set up JMeter to send SQL queries to the database and collect the results for analysis.

Procedure:

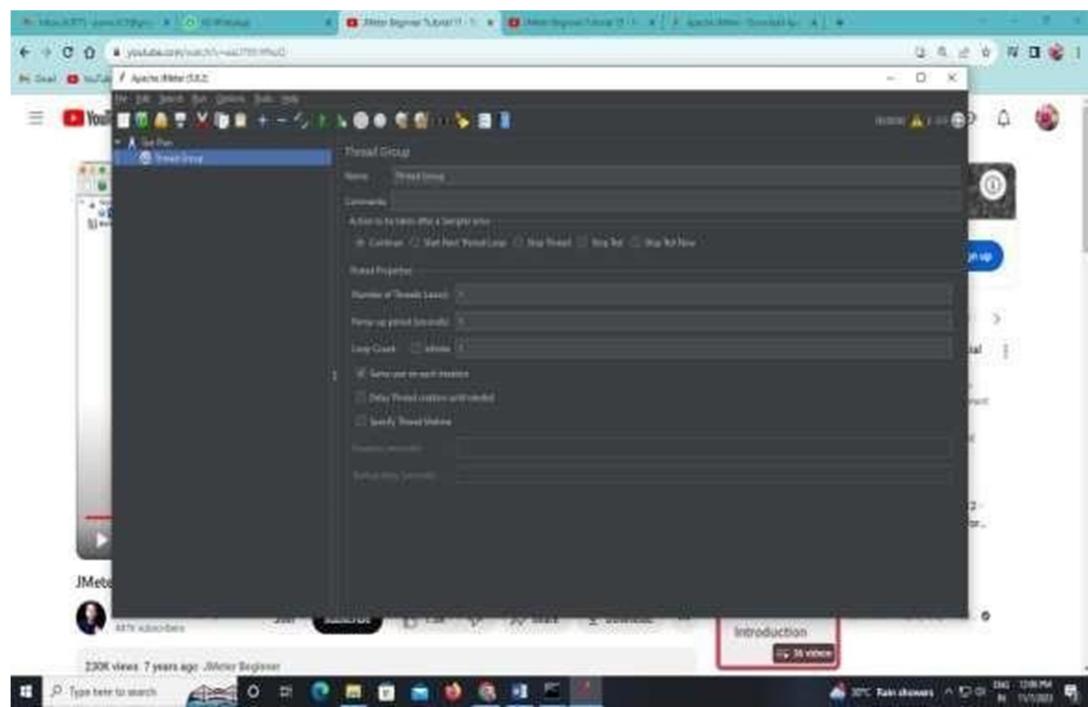
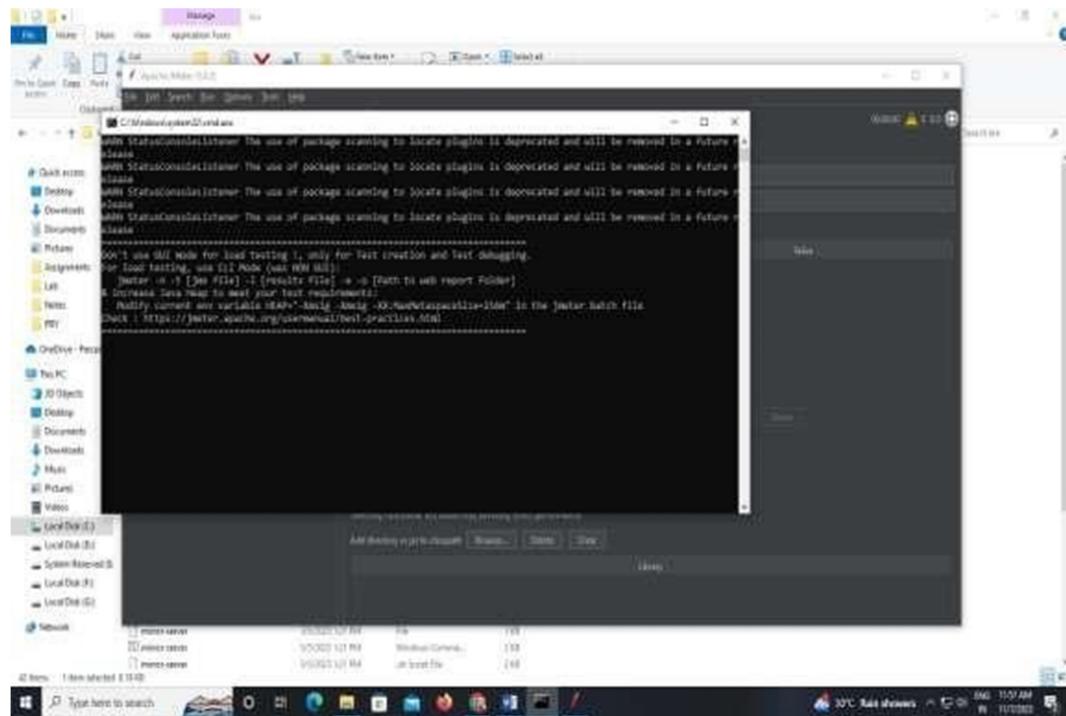
1. Install Required Software:

- **Install JMeter:** Download and install JMeter from the official Apache JMeter website.
- **Install SQL Server and SQL Management Studio:** If you haven't already, set up SQL Server and SQL Management Studio to manage your database.



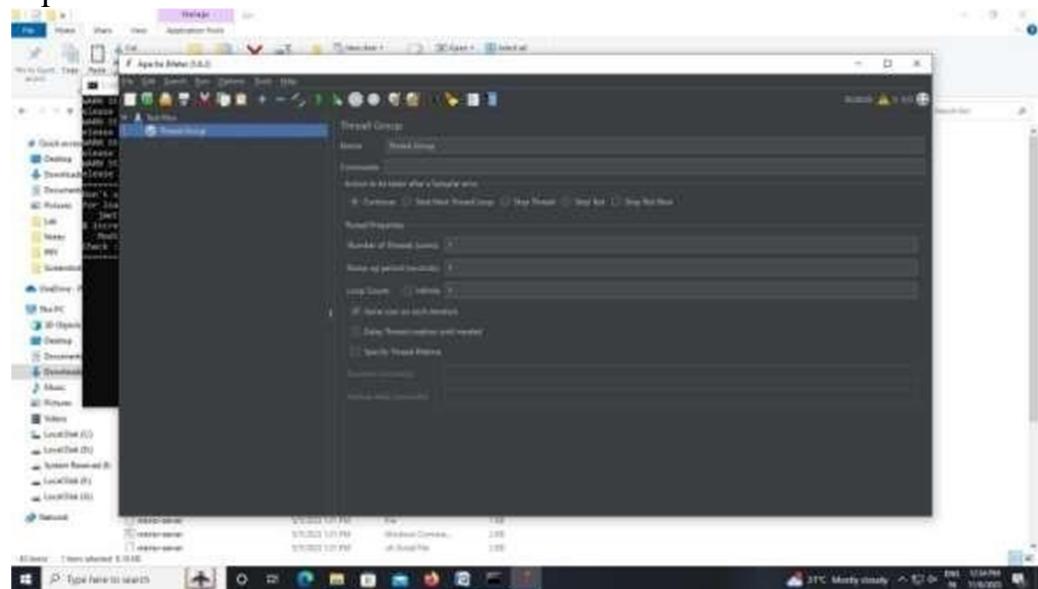
2. Create a Test Plan in JMeter:

- Launch JMeter and create a new Test Plan.



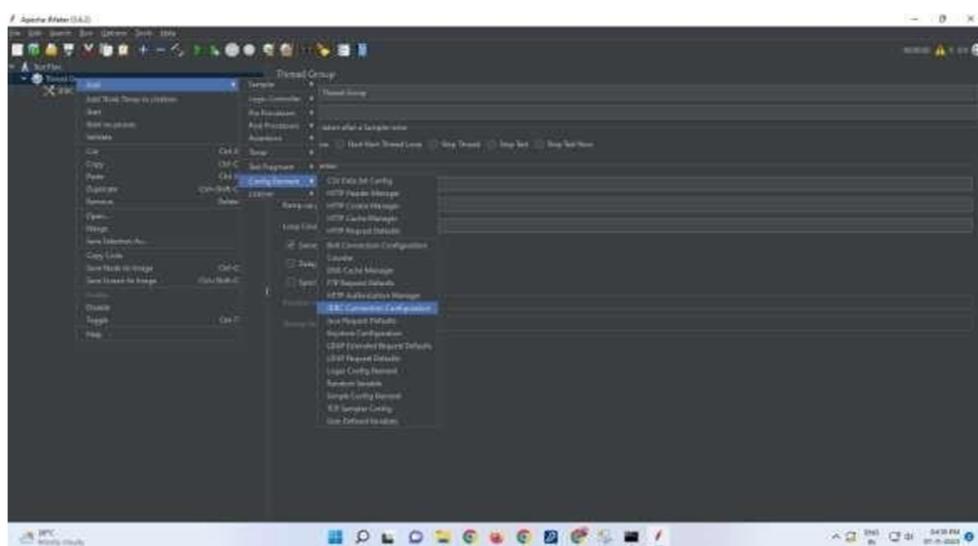
3. Add Thread Group:

- Add a Thread Group to your Test Plan to simulate the number of users and requests.



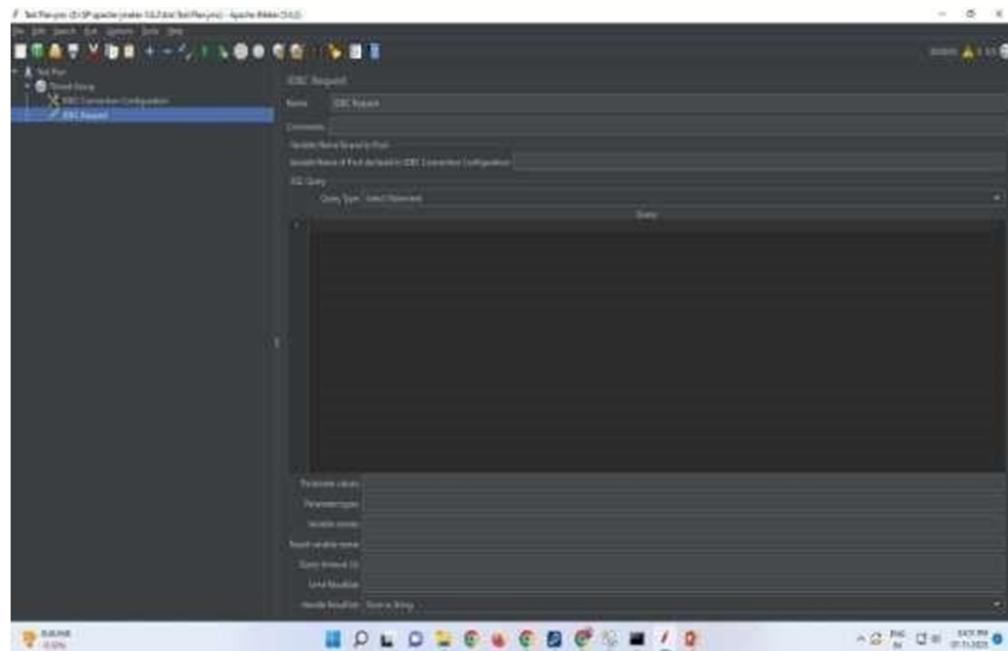
4. Add JDBC Connection Configuration:

- Add a JDBC Connection Configuration element to your Thread Group. Configure it with the database connection details, such as the JDBC URL, username, and password. This element will allow JMeter to connect to your SQL Server database.



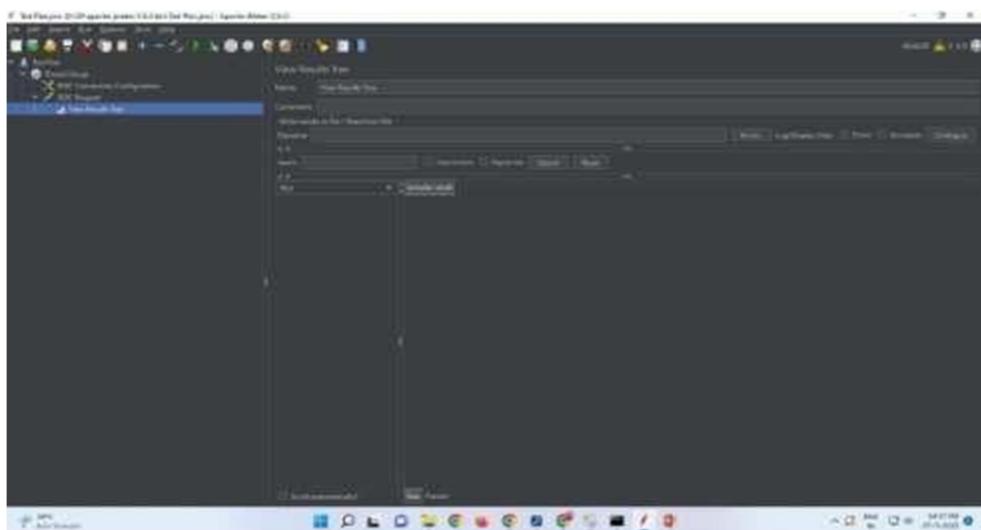
5. Add a JDBC Request Sampler:

- Add a JDBC Request sampler to your Thread Group. This sampler will contain your SQL query.
- Configure the JDBC Request sampler with the JDBC Connection Configuration created in the previous step.
- Enter your SQL query in the "Query" field of the JDBC Request sampler.



6. Add Listeners:

- Add listeners to your Test Plan to collect and view the test results. Common listeners include View Results Tree, Summary Report, and Response Times Over Time.

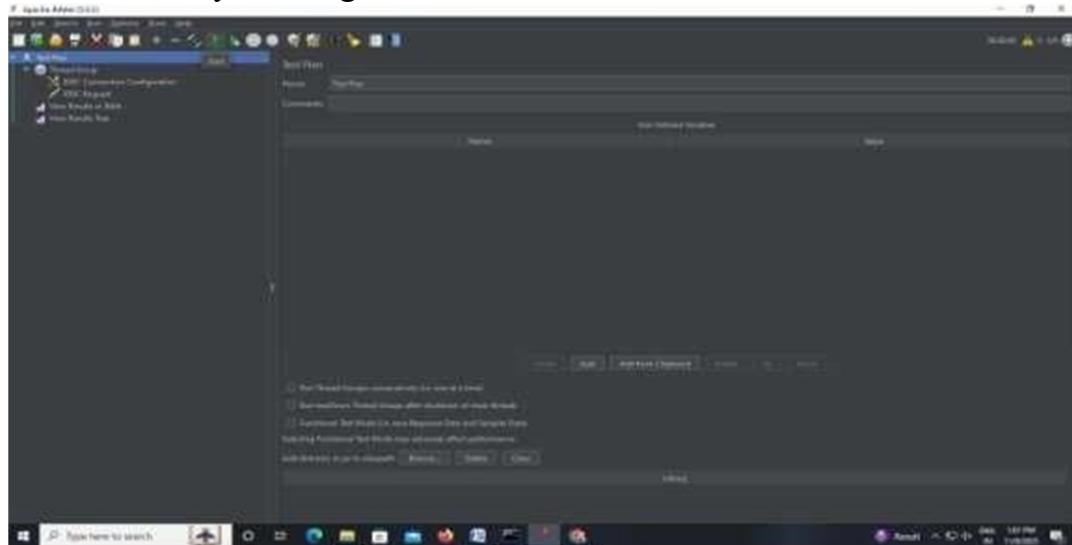


7. Configure Your Test Plan:

- Configure the number of threads (virtual users), ramp-up time, and loop count in the Thread Group to simulate the desired load.

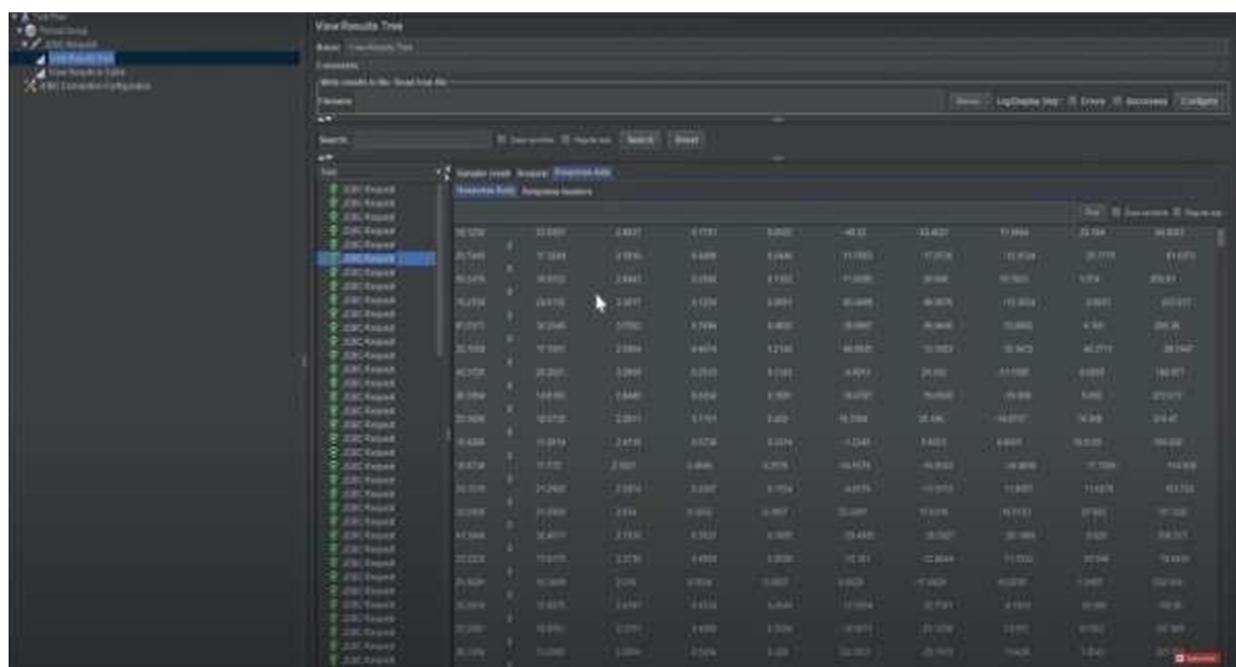
8. Run the Test:

- Start the test by clicking the "Run" button in JMeter.



9. View and Analyse Results:

- After the test has completed, you can view and analyse the results using the listeners you added. You can analyse response times, errors, and other performance metrics.



10. Optimize and Fine-Tune:

- Based on the results, you can optimize your SQL queries and JMeter test plan to finetune the performance of your database.

Conclusion

Using JMeter in conjunction with SQL Management Studio can be a powerful combination for load testing and performance analysis of applications that rely on SQL Server databases. This approach allows you to simulate a realistic user load, send SQL queries to the database, and evaluate the system's performance under various conditions.

JMeter in combination with SQL Management Studio provides a robust solution for assessing the performance of applications that rely on SQL Server databases. Through thorough testing, analysis, and optimization, you can ensure your application is capable of delivering a reliable and responsive experience to users even under heavy load conditions.